

IHP SG13G2 · OPEN SOURCE RFIC EM WORKFLOW

Using openEMS with IHP SG13

User's Guide

Volker Mühlhaus
volker@muehlhaus.com

Document version 2026-09-02 (v3)

Contents

1	What's New	3
2	About this workflow	3
2.1	Method of Moments / FEM vs. FDTD	3
3	Getting the code: pip package vs. local copy	4
3.1	Recommended: install gds2openEMS as a Python module	4
3.2	Alternative, no longer recommended: local modules copy	5
4	Required software and Python modules	5
5	Model configuration syntax: settings{} dict vs. loose variables	5
6	Quick tour	5
6.1	Minimum configuration	6
6.2	Input files	6
6.3	Port definitions	6
6.4	Meshing	8
6.5	Running the model: preview, mesh, simulate	9
6.6	Re-running a model: the hash-based skip	10
6.7	Output files and Touchstone conversion	11
7	The settings{} dictionary in detail	11
7.1	Required settings	12
7.2	Optional settings	13
8	Simulation volume and boundaries	13
9	Field dumps	15
10	Antenna simulation	15
11	Port parasitics and de-embedding	17
12	Advanced topics	17
12.1	Forcing the solver thread count	17
12.2	Automatic meshing with easyMesh4openEMS	17
12.3	Overriding XML stackup <Variable> from the model code	17
13	Examples	18
14	FAQ	18
15	Appendix	20
15.1	Keeping the local-copy example in sync	20
15.2	List of examples	20

1 What's New

This chapter gives a brief overview of major changes since the previous edition (v2, March 2025) of this guide. For the complete, dated change log, see [CHANGES.md](#) in the repository.

- One major change is that documentation and examples now **install gds2openEMS as a module from PyPi** (`pip install gds2openEMS`), instead of the local copy of the modules folder used before. The old method with local copy still works, but the new method using Python module installation is more convenient.
- The second major change is that examples and documentation now use the **settings[] dictionary syntax for simulation model settings**, similar to the gds2palace workflow. The previous model syntax with loose variables is still supported, so you can choose whatever suits your needs. Long term, the settings[] dictionary syntax is preferred because it allows to switch between openEMS and Palace workflows more easily.
- A graphical **XML stackup editor** was added in the `stackup_editor` directory. When the gds2openEMS Python module is installed, you can call this from the commandline using `stackupEditor`, with an XML file as optional parameter. Because of this GUI-based stackup editor, the PySide6 module is now a required dependency, and installed automatically when you install the gds2openEMS module.
- **Derived layers and <Variable> overrides** in the XML stackup format, see [Overriding XML stackup <Variable> from the model code](#) and [XML_stackup_format.md](#).
- The `postprocess_only` setting is no longer required. `runSimulation()` already skips the FDTD solve automatically via a content hash of the model. If nothing changed since the last run, it reuses the existing result instead of re-simulating.
- **shapely is now a hard dependency**. It is used for splitting keyhole and hole polygons before CSXCAD extrusion. Layouts with cutouts no longer need a `preprocess_gds` flag at all.
- Port de-embedding script `scripts/deembed_openEMS.py` is a “quick & dirty” solution to estimate and remove lumped port’s parasitic inductance, see [Port parasitics and de-embedding](#).
- `numThreads` setting to **force the solver thread count** to a fixed value, instead of relying on automatic detection.

2 About this workflow

openEMS is an open-source EM simulation tool based on the FDTD (Finite Difference Time Domain) method. This workflow generates openEMS simulation models directly from **GDSII layout files** and an **XML technology stackup**, for RFIC structures on the IHP SG13G2 process — inductors, transmission lines, couplers, antennas and similar passive structures, producing S-parameters.

Two files must be provided by the user: the **layout in GDSII format**, and a **simulation model script** in Python. The model script references an XML stackup file and calls functions from the workflow code to build the openEMS/CSXCAD model, mesh it, run the FDTD solve, and post-process the result into Touchstone S-parameters. That workflow code can be the gds2openEMS package, or a local `modules/` copy — see the next chapter.

2.1 Method of Moments / FEM vs. FDTD


It’s useful to understand how FDTD’s simulation flow differs from a frequency-domain solver. Frequency-domain solvers include Method of Moments, as used in Momentum or Sonnet, and FEM, as used in the sibling `gds2palace_ihp_sg13g2` workflow:

- **MoM/FEM** is a frequency-domain solver: it solves once per frequency point and gets the **full S-matrix** for that one frequency. Covering a frequency range means repeating the solve at every frequency of interest.

- **FDTD** is a time-domain solver: exciting one port with a wideband Gaussian pulse and running the field propagation forward in time gives, after an FFT, **one column of the S-matrix across the entire frequency range** in a single run. To get the full S-matrix, you repeat the run with excitation at each other port in turn.

Method of Moments, FEM

- Frequency domain solver
- Full S-Matrix for one frequency



FDTD

- Time domain solver, pulse excitation
- One column of S-matrix for complete frequency range



FFT time domain -> frequency domain




Figure 1: Comparison between MoM/FEM and FDTD simulation methods

The number of frequency points you request has essentially no effect on FDTD simulation time — the frequency-domain result is FFT post-processing of the time-domain data, so there's no speed penalty to specifying `numfreq = 401` instead of a coarser sweep.

3 Getting the code: pip package vs. local copy

There are two independent ways to get this workflow's code into your model script — pick one, they're interchangeable, and this choice is separate from the `settings{}` vs. `loose-variable` syntax choice below.

3.1 Recommended: install gds2openEMS as a Python module

```
python3 -m venv ~/venv/openems
source ~/venv/openems/bin/activate
pip install gds2openEMS
```

Your model script then does:

```
from gds2openEMS import *
```

and needs no local copy of the workflow code at all — just your GDSII layout, your XML stackup file, and the model script itself. Update with `pip install gds2openEMS --upgrade`. All examples under `workflow/` use this.

3.2 Alternative, no longer recommended: local modules copy

Clone/download this repository and copy the workflow/modules/ folder next to your own model script, which then does:

```
import sys, os
sys.path.insert(0, os.path.abspath(os.path.join(os.path.dirname(__file__), 'modules')))
from modules import *
```

This is still fully supported — useful if you want to vendor a specific version, work offline, or edit the workflow code itself. `more_examples/local_modules_copy/` is a fully self-contained example written this way: its own modules/ copy, GDSII layout and XML stackup, no PyPI package needed.

4 Required software and Python modules

This workflow is based on the Python bindings for openEMS; please refer to <https://www.openems.de/> and <https://docs.openems.de/python/install.html#python-linux-install> for building/installing openEMS and CSXCAD themselves. Pre-built binaries are available for Windows, which is the easier path if you don't want to build from source; for Linux build issues, see the openEMS forum at <https://github.com/thliebig/openEMS-Project/discussions>.

In addition to openEMS/CSXCAD, these Python modules must be installed: `gdspy`, `shapely`, `numpy`, `matplotlib`. `pip install gds2openEMS` declares `gdspy`, `CSXCAD`, `openEMS`, `numpy` and `shapely` as its own dependencies, and `pip` will attempt to pull them in. Note that openEMS/CSXCAD's native bindings are normally built from source rather than installed as an ordinary wheel — see the openEMS install docs linked above. `matplotlib` is only needed by the example scripts' own plotting code, not by the installed package itself.

5 Model configuration syntax: settings{} dict vs. loose variables

Model scripts pass their configuration to `setupSimulation()`/`runSimulation()` two ways — pick one, the underlying functions accept both, and you can mix either syntax with either distribution method above:

- `settings = {}` **dictionary** is the recommended style. Every option is a key in one dict, passed as `setupSimulation(FDTD=FDTD, settings=settings)`. This is the style all workflow/ examples use now, and it's also what `gds2palace_ihp_sg13g2` (the sibling AWS Palace/Elmer FEM workflow) uses. As a bonus, `max_cellsize` is computed for you from `fstop/unit/cells_per_wavelength` instead of you having to compute it yourself.
- **Loose top-level variables** (`fstart = 0e9`, `refined_cellsize = 1`, ...) passed individually and positionally — the original style, still fully supported. `more_examples/local_modules_copy/run_line_viaport.py` is a current example of it.

The rest of this guide shows the `settings{}` style, since that's what all current workflow/ examples use; anywhere it says `settings['xyz'] = ...`, the loose-variable style has an equivalent plain `xyz = ...` passed positionally instead.

6 Quick tour

This chapter gives a brief, practical walkthrough. It uses `workflow/run_inductor_diffport.py` (a 1-port octagon inductor) as the running example.

6.1 Minimum configuration

If you installed gds2openEMS via pip, the minimum configuration for a model is just three files in one directory: the **XML stackup**, the **GDSII layout**, and your **model script**. No modules/ folder needed.

If you're using the local-copy method instead, add a fourth item: your own copy of the modules/ folder, next to the script.

6.2 Input files

```
gds_filename = "L_2n0_simplified.gds" # geometries
XML_filename = "SG13G2.xml"          # stackup

settings['preprocess_gds'] = False    # see note below
settings['merge_polygon_size'] = 1.0  # merge via arrays within this distance (microns), 0 to disable

materials_list, dielectrics_list, metals_list = stackup_reader.read_substrate(XML_filename)
layernumbers = metals_list.getlayernumbers()
layernumbers.extend(simulation_ports.portlayers)

allpolygons = gds_reader.read_gds(gds_filename, layernumbers, purposelist=[0],
                                   metals_list=metals_list,
                                   preprocess=settings['preprocess_gds'],
                                   merge_polygon_size=settings['merge_polygon_size'])
```

preprocess_gds is a legacy setting kept only for backward compatibility with older model scripts. The GDSII reader now always splits keyhole/hole polygons safely after flattening the layout, so this option is a no-op today. Setting it to True just prints a note that it's obsolete. New model scripts don't need to set it at all.

merge_polygon_size applies to layers declared Type="via" in the XML stackup. With a non-zero value, polygons on that layer are oversized by half that distance, overlapping ones are merged, then undersized back — merging a via array into one bounding-box shape, as long as the maximum via spacing doesn't exceed the given value.

Merging fills the gaps between the vias with via material, so the merged block conducts better than the real via array. The via conductivity in the SG13G2 stackup files is derived from the per-via resistance in the process specification, i.e. it is valid for the via cross section only. With settings['fill_factor_correction'] = True, setupSimulation() corrects this: each merged via polygon gets its conductivity multiplied by its fill factor, the original via area inside it divided by the merged area (rounded to 2 decimals). Polygons with different fill factors get separate materials, named like TopVia2_x0.49, which you can see in AppCSXCAD; single vias that were not merged keep fill factor 1 and the original material. The correction needs merge_polygon_size > 0, because the fill factor is calculated from the via polygons before merging: a GDSII file where via arrays were already merged into blocks (e.g. by gds_prepare_for_EM) gets no correction. PEC via layers are not scaled.

6.3 Port definitions

Ports are created from polygons on special GDSII layers (by convention, layer 201 and above), not by coding their position directly in Python:

```
simulation_ports = simulation_setup.all_simulation_ports()

# via port: from_layername / to_layername, direction is the vertical axis
simulation_ports.add_port(simulation_setup.simulation_port(
    portnumber=1,
    voltage=1,
    port_Z0=50,
    source_layernum=201,
    from_layername='Metal1',
    to_layername='TopMetal2',
    direction='z'))
```

```
# in-plane port: target_layername, direction is the propagation axis in-plane
simulation_ports.add_port(simulation_setup.simulation_port(
    portnumber=1,
    voltage=1,
    port_Z0=50,
    source_layernum=201,
    target_layername='TopMetal1',
    direction='x'))
```




Figure 2: Port layout defined in GDSII file on layer 201

Each port must use its own GDSII source layer. `voltage=0` disables that port's own excitation without removing its definition.

In scripts that build their excitation list from `simulation_ports.all_active_excitations()` — like `run_generic_nport.py` — this also skips that port's FDTD run entirely, saving simulation time. Scripts with a hardcoded excitation list (most of the workflow/ examples) must leave that port's number out of the list themselves; `voltage=0` alone does not skip anything there, since the script controls which ports run, independent of their voltage value.

A skipped port's S-parameters are simply not computed. If you then try to write a Touchstone file that references one of them, `write_snp()` raises a clear error naming the missing S-parameters, rather than producing an incomplete or corrupt file. See the FAQ.

The port can be mapped to a metal layer as shown here, with x or y orientation, but it can also be mapped to a via layer and used in z direction, if necessary. This gives maximum flexibility in creating EM ports.

The port direction also supports reverse polarity, by adding a minus sign to the direction. This can be required to place the reference terminal on the correct side, as shown in the example below.




Figure 3: Example for polarity in port direction

Composite ports combine multiple EM ports into one logical port, for example a GSG pad pair. These are not built into `simulation_setup`'s port model directly.

The GSG port on each end of the line consists of 2 port definitions each, with one port in y direction and the other port in -y direction, to have the same effective polarity. And because they are effectively in parallel, we specify $2 \cdot Z_0$ as port impedance, to have Z_0 in total.




Figure 4: Port layout for GSG configuration, composite port from 2 port elements

`run_line_GSG_complex.py` shows the manual approach: simulate the underlying EM ports together, then combine their results in your own post-processing code.

6.4 Meshing

The default in all examples is **automatic meshing based on geometry**: mesh lines are placed to follow polygon edges and diagonal features. Lines that end up too close together are then merged or removed, since they would slow simulation without adding accuracy. `settings['refined_cellsize']` sets the target mesh resolution at conductor edges. `settings['cells_per_wavelength']` sets the coarse mesh size elsewhere in the model: the maximum cell size is the wavelength at `fstop` in the densest dielectric, divided by `cells_per_wavelength`.




Figure 5: Automatic meshing

In z direction, mesh lines come from the stackup. Every dielectric boundary and the bottom and top of every metal and via layer used in the layout get a mesh line, and lines closer than 1 nm are merged into one. Metals are then subdivided:

- a metal up to $3 \times \text{refined_cellsize}$ thick is divided into cells of about refined_cellsize , split at any dielectric interface inside it (e.g. a passivation that cuts TopMetal2 halfway), so no subdivision line lands next to such an interface;
- a thicker metal gets one extra line at refined_cellsize from its bottom and top surface.

Dielectrics are graded between these lines, up to the maximum cell size. So the smallest z cell is set by the thinnest layer actually used, not by the meshing rules.

Before 28-Sep-2026, the z mesh placed metal subdivision lines without regard to interfaces inside the metal, which created very small cells in some stackups (0.1 μm in the L6n2 passivation-cut stackup at 1 μm refined_cellsize). To reproduce results from that version, set `settings['z_mesh_function'] = util_meshlines.create_z_mesh_legacy`.

The default meshing method described above is reliable and easy to work with, and is a good choice for most models. The optional example `more_examples/easyMesh/` demonstrates an alternative meshing engine, shown for information — in most cases the two approaches are equivalent, and neither is a general replacement for the other. See chapter [Automatic meshing with easyMesh4openEMS](#).

6.5 Running the model: preview, mesh, simulate

```
FDTD = simulation_setup.setupSimulation(FDTD=FDTD, settings=settings)
subl_data_path = simulation_setup.runSimulation(FDTD=FDTD, settings=settings)
```

With `settings['preview_only'] = True` and `settings['no_gui']` left at its default (False), running the script pops up **AppCSXCAD**, the 3D viewer, showing the meshed model before any simulation starts:




Figure 6: AppCSXCAD 3D mesh view of a differential-port octagon inductor

To close the preview, just close that window. You can also press Ctrl-C to abort. No other user action is required. Once you're happy with the model and mesh, set `settings['preview_only'] = False` and re-run. The FDTD solve then starts as soon as you close the preview window. If `settings['no_gui'] = True`, the solve starts immediately, since there is no preview window to close.

6.6 Re-running a model: the hash-based skip

`runSimulation()` computes a content hash covering both the generated CSX model file and the calling model script (.py) itself, up to and including the line that calls `runSimulation()`. Code further down in that script (post-processing, plotting, ...) is not part of the hash. It compares that hash against the hash stored from the previous run in that model's output directory. If the hashes match, it skips the FDTD solve and prints:

```
Data for this model already exists, skipping simulation!
To force re-simulation, add parameter "force_simulation=True" to the runSimulation() call.
```

To force a re-solve even when the hashes match, pass `force_simulation=True` to `runSimulation()`.

The AppCSXCAD preview itself is **not** gated by this hash check. It is shown on every run, for the first port excitation, regardless of whether the model changed. Set `no_gui = True` to suppress it. This way you always get a chance to look at the current model before deciding whether to let it re-simulate. Only the FDTD solve is skipped when nothing changed. So you can freely re-run a script after only changing your own post-processing or plotting code, without waiting through a re-solve, while still seeing the preview every time.

6.7 Output files and Touchstone conversion

Output goes to `output/<model_basename>_data/sub-<n>/` for each port excitation `n`. This folder contains the CSX/mesh XML and the time-domain port voltage/current data. Field dumps are written there too, if you requested them.

There is no separate conversion tool for openEMS. Every example model script computes the S-parameters itself after the FDTD solve finishes, and calls `utilities.write_snp()` to write them straight to a Touchstone `.snp` file:

```
snp_name = os.path.join(sim_path, model_basename + '.s' + str(num_ports) + 'p')
utilities.write_snp(s_params, f, snp_name, z0=simulation_ports.get_reference_impedance())
```

So the `.snp` file already exists once your model script finishes running. Just run the script:

```
python your_model.py          # generates the model, runs the simulation, writes the .snp file
```

7 The settings{} dictionary in detail

Settings keys are case-insensitive: `settings['numthreads']` works like `settings['numThreads']`, and `setupSimulation()` prints a note when it uses a key in a different case. If both spellings are set with different values, the spelling shown in the tables below wins, with a warning.

7.1 Required settings

Key	Meaning
<code>settings['unit']</code>	Unit of geometry values, typically 1e-6 (microns)
<code>settings['margin']</code>	Oversize of dielectrics from the GDSII bounding box, in the geometry unit
<code>settings['fstart'] / settings['fstop']</code>	Start/stop frequency in Hz for the S-parameter output (FFT post-processing — see About this workflow)
<code>settings['numfreq']</code>	Number of frequency points in the output sweep
<code>settings['refined_cellsize']</code>	Target mesh size at conductor edges
<code>settings['Boundaries']</code>	Required, no built-in default. List of 6 boundary conditions, one per side (xmin, xmax, ymin, ymax, zmin, zmax): 'PEC' (lossless metal box), 'PMC' (magnetic wall, useful for symmetry), 'MUR' (simple absorbing), or 'PML_8' (higher-quality absorbing, much slower simulation)
<code>settings['energy_limit']</code>	Residual energy (dB) at which the FDTD time-domain solve is considered converged
<code>settings['cells_per_wavelength']</code>	Coarse-mesh resolution away from refined edges, 10 or more (the examples use 10-20). The maximum cell size is calculated from it as $\text{wavelength at } f_{\text{stop}} / (\sqrt{\epsilon_{r,\text{max}}} \cdot \text{cells_per_wavelength})$. Alternatively, leave out <code>cells_per_wavelength</code> and set <code>settings['max_cellsize']</code> directly, in the geometry unit

7.2 Optional settings

Key	Default	Meaning
<code>settings['preview_only']</code>	False	Show the AppCSXCAD preview and stop, without simulating
<code>settings['no_gui']</code>	False	Never show AppCSXCAD; always proceed straight to simulation once the model changed (see the hash-based skip)
<code>settings['force_simulation']</code>	False	Re-simulate even if the model hash matches a previous run
<code>settings['max_cellsize']</code>	calculated	Maximum mesh cell size in the geometry unit. Only used when <code>fstop</code> , <code>unit</code> or <code>cells_per_wavelength</code> is missing; otherwise the value calculated from <code>cells_per_wavelength</code> is used
<code>settings['preprocess_gds']</code>	False	Legacy, now a no-op — see Input files
<code>settings['merge_polygon_size']</code>	0	Merge via-array polygons within this distance — see Input files
<code>settings['fill_factor_correction']</code>	False	Scale the conductivity of merged via arrays by their via fill factor — see Input files
<code>settings['air_around']</code>	0	Extra air spacing around the model in addition to margin, single value or a 6-element list [<code>xmin</code> , <code>xmax</code> , <code>ymin</code> , <code>ymax</code> , <code>zmin</code> , <code>zmax</code>]
<code>settings['numThreads']</code>	automatic	Force the openEMS solver thread count — see Forcing the solver thread count
<code>settings['easyMesh']</code>	False	Use the easyMesh4openEMS automatic mesh generator instead of the built-in one — see Automatic meshing with easyMesh4openEMS
<code>settings['field_dumps']</code>	none	A <code>simulation_setup.all_field_dumps()</code> object — see Field dumps

8 Simulation volume and boundaries

The simulation volume is the GDSII geometry bounding box, oversized by margin on all sides. If `air_around` is set, that adds further oversize. The volume is then meshed into cells according to `refined_cellsize` and `cells_per_wavelength`.




Figure 7: Margin defined in simulation settings

In the model code, boundaries are listed in this order: $x_{\min}$, $x_{\max}$, $y_{\min}$, $y_{\max}$, $z_{\min}$, $z_{\max}$. Each of the six outer boundaries is independently metal (PEC), magnetic-wall (PMC), or absorbing. Absorbing boundaries are MUR, or the more accurate but more expensive PML_8. See `settings[ 'Boundaries' ]` above. A PMC boundary is useful to exploit a structure's symmetry and simulate only half the model. An absorbing boundary is needed wherever radiation should leave the simulation domain rather than reflect back in — for example all six sides for an antenna, see [Antenna simulation](#).

Metal layers are often embedded inside a dielectric, so two materials would otherwise overlap. The workflow resolves this by construction: dielectric layers are the base fill, and metal/via geometry is meshed as its own conductor sheets or volumes on top, matching the Z-ranges defined for each layer in the XML stackup file.

9 Field dumps

For visualization in AppCSXCAD or ParaView, you can request field data to be written during simulation:

```
field_dumps = simulation_setup.all_field_dumps()
field_dumps.add_frequency_dump(name='Jf',
    file_type='vtk',
    dump_type='J',
    frequency=30e9,
    source_layernum=301,
    from_layernum='TopMetal2',
    to_layernum='TopMetal2',
    offset_top=0,
    offset_bottom=0)
field_dumps.add_time_dump(name='Et',
    file_type='vtk',
    dump_type='E',
    source_layernum=302,
    from_layernum='Metal1',
    to_layernum='TopMetal2',
    offset_top=10,
    offset_bottom=0)
settings['field_dumps'] = field_dumps
```

Each dump is positioned by its own GDSII bounding-box layer (like a port), with dump_type one of 'E', 'H', 'J', 'rotH', written as 'vtk' or 'hdf5'. That layer must be read from the GDSII file, like the port layers: add `layernumbers.extend(field_dumps.dumplayers)` before `read_gds()`. Otherwise the dump silently covers the bounding box of the entire layout, and `setupSimulation()` prints a warning. The dump box is the bounding rectangle of all polygons on that layer. Its edges also become mesh lines, like any other polygon edge. A frequency dump captures the steady-state field at one frequency; a time dump captures the field evolving during the FDTD solve. See `workflow/run_line_viaport_fielddump.py` for a complete example.

10 Antenna simulation

`workflow/run_dual_dipole.py` demonstrates far-field antenna pattern calculation (based on a 245 GHz dual-dipole design by IHP's Klaus Schmalz et al.). Antenna models differ from the passive-structure examples in three ways:

- All six outer boundaries are absorbing, either MUR or PML_8, since radiated power must leave the simulation domain rather than reflect back. Using an 'MUR' boundary with a somewhat larger `air_around/margin` is usually faster than a tighter 'PML_8' box for antenna problems. PML boundaries slow down openEMS a lot!
- In the simulation setup for this antenna, the optional parameter `air_around` adds additional air spacing all around the model, to ensure proper distance from the absorbing boundaries.
- An **NF2FF (near-field to far-field) box** is added around the radiator, which openEMS uses after the time-domain solve to compute the radiation pattern:

```
nf2ff_box = FDTD.CreateNF2FFBox(opt_resolution=[max_cellsize]*3, frequency=[ftarget])
...
nf2ff_res = nf2ff_box.CalcNF2FF(subl_data_path, ftarget, theta, phi)
```




Figure 8: AppCSXCAD mesh view of the dual-dipole antenna model

From the NF2FF result and the port's accepted/incident power (`CSXport.P_acc/P_inc`), the model script computes total radiated power, peak directivity, radiation efficiency and mismatch loss, and plots directivity/gain vs. angle. See `run_dual_dipole.py`'s own code for the full calculation.




Figure 9: Return loss and directivity

This model needs `max_cellsize` for the NF2FF box resolution, and that value is needed *after* calling `setupSimulation()`. So it computes `wavelength_air/max_cellsize` manually, rather than relying on `setupSimulation()`'s internal settings-dict auto-computation. That computed value isn't written back into settings after the call.

The call to function `CreateNF2FFBox()` is plain openEMS syntax, so check out the documentation available there:

- <https://docs.openems.de/python/openEMS/nf2ff.html>

- <https://github.com/thliebig/openEMS/blob/master/python/openEMS/nf2ff.py>

Parameter `frequency` is optional, to store only the field data for that one frequency, which saves disk space. Without this optional parameter, `nf2ff` stores all time domain data, and patterns can be calculated in postprocessing for any frequency. The disadvantage is that data files will be much larger then.

For the evaluation of the NF2FF data files and creating the antenna pattern plot, we also use plain openEMS syntax. Have a look at example `run_dual_dipole.py` for more details.

11 Port parasitics and de-embedding

This workflow creates **lumped ports**, which introduce some physical length/area into the model and therefore some parasitic series inductance — port size that isn't small compared to the device under test will show up as extra inductance in the raw S-parameter result.

`scripts/deembed_openEMS.py` is a “quick & dirty” port de-embedding script: it reads the `port_information.json` geometry metadata that `gds2openEMS` writes automatically for each run, estimates the parasitic inductance built into each lumped port using a thin-sheet approximation, and removes it per port by cascading a negative-inductance section:

```
python3 deembed_openEMS.py inputfile.s2p
```

Output is written to a new file with suffix `_deembedded`. This works for any port count, but the method is approximate and **not applicable to composite port configurations** like `run_line_GSG_complex.py`'s GSG ports — treat it as a useful estimate, not a precise correction.

12 Advanced topics

12.1 Forcing the solver thread count

By default, openEMS auto-detects the “best” thread count, which can end up low if the machine is busy with other work. `settings['numThreads']` overrides that:

```
settings['numThreads'] = 8 # 0 (or omitted) = automatic detection
```

See `more_examples/numThreads/` for complete examples in both the `settings{}` and loose-variable styles.

12.2 Automatic meshing with easyMesh4openEMS

`more_examples/easyMesh/openEMS_generic_nport_MA.py` demonstrates `easyMesh4openEMS`, an alternative automatic mesh-line placement engine, enabled via:

```
settings['easyMesh'] = True
```

`easyMesh4openEMS` is a separate package (`pip install easyMesh4openEMS`) not installed by default alongside `gds2openEMS`. If it isn't installed, `settings['easyMesh'] = True` doesn't crash the run — `setupSimulation()` catches the missing import, prints a note that it's being ignored until you install the package, and falls back to the default geometry-based automatic meshing.

12.3 Overriding XML stackup <variable> from the model code

In the early days of this workflow, any change to the technology stackup like thickness of bulk silicon or thickness of air layer above required a separate, modified XML file.

To simplify this, an extension to the stackup file format was introduced with variable in the XML stackup file. You can now override these variables from your model code, to control values like thickness without using another XML file.

`more_examples/parameterized_XML_stackup/` shows how to override stackup variables `total_thickness`, `air_thickness` from the model script instead of editing the XML file directly. This is done in the `read_substrate()` code line:

```
materials_list, dielectrics_list, metals_list = stackup_reader.read_substrate(
    XML_filename, variable_overrides={'total_thickness': 100, 'air_thickness': 100})
```

Two versions of the same model are provided — `run_rfcmm_2port_parameters.py` (loose variables) and `run_rfcmm_2port_parameters_settings.py` (`settings{}` dict) — producing an identical model, to show both syntax options side by side. See the XML stackup's `<Variables>` block and `"="`-prefixed expressions in [XML_stackup_format.md](#) for the full format.

13 Examples

All examples read GDSII + XML stackup and write a Touchstone S-parameter file (except where noted). For a longer description of each, see the top-level [README.md](#).

workflow/ (primary examples, `settings{}` + `gds2openEMS`):

- `run_line_viaport.py` — thru line with via ports, single-side excitation with symmetry assumed for the reverse path
- `run_line_full2port.py` — thru line, full 2-port excitation
- `run_line_GSG_complex.py` — thru line with composite GSG ports
- `run_line_noGDSII.py` — geometry built entirely from Python code, no GDSII input
- `run_line_viaport_fielddump.py` — via-port line with field dumps (see [Field dumps](#))
- `run_inductor_diffport.py` — 1-port octagon inductor (differential mode)
- `run_inductor_2port.py` — 2-port version of the same inductor, via ports to a common ground
- `run_generic_nport.py` — generic n-port model, excites every defined port in turn for a full S-matrix
- `run_rfcmm_2port_full.py` — MIM capacitor with via-array merging
- `run_dual_dipole.py` — antenna example, see [Antenna simulation](#)

more_examples/:

- `local_modules_copy/run_line_viaport.py` — the local-copy + loose-variable distribution/syntax combination, self-contained (own modules/, GDS, XML)
- `parameterized_XML_stackup/` — XML `<Variable>` overrides, both syntax styles
- `numThreads/` — forcing solver thread count, both syntax styles
- `easyMesh/openEMS_generic_nport_MA.py` — easyMesh4openEMS automatic meshing
- `resistors_sg13g2/run_resistors_Rsil.py` — IHP resistor derived-layer support
- `combine_layout_sources/run_combine_example.py` — combining GDSII-sourced and manually-coded geometry (rectangles/polygons added directly in Python) in one model

14 FAQ

How do I use the example models? Copy the example script, its GDSII file and XML stackup into a new directory, or into your own project directory. If you're using `gds2openEMS` via pip, no `modules/` copy is needed there. Then modify the configuration to match your own layout and requirements.

How do I create my own model? Start from the example that's closest to your use case, rather than starting from scratch. Match the port type, via or in-plane, and the port count. Then adapt the GDSII/XML filenames, port definitions and settings{} values.

Does voltage=0 actually skip that port's simulation? It depends on how the script builds its excitation list. `run_generic_nport.py` and a few other examples call `simulation_ports.all_active_excitations()` to get the list of ports to excite, which already filters out `voltage=0` ports — for those scripts, a `voltage=0` port's FDTD run is skipped entirely. Most workflow/ examples instead hardcode which port numbers to excite (e.g. `[1]` or `[[1], [2]]`); there, `voltage=0` by itself changes nothing — you must also remove that port number from the hardcoded list yourself. Either way, a port that was never excited has no S-parameters computed for it; trying to write a Touchstone file that references one raises a clear error naming the missing S-parameter(s), rather than producing a corrupt file.

What is a good mesh size? There's no single answer — it depends on the smallest feature (gap or line width) you need to resolve accurately, and on how much time you can afford. As a starting point, a `refined_cellsize` around 1/5 to 1/10 of your smallest critical dimension is reasonable; check convergence by comparing results at two different `refined_cellsize` values before trusting the finer one.

“Warning: Unused primitive (type: LinPoly) detected in property ...”: These messages pop up when openEMS solver detects a layout with multiple touching metals on the same layer and same priority. It is purely cosmetic in this workflow and can be safely ignored.

Why is my simulation slow? FDTD simulation time scales with mesh cell count and the number of time steps needed to reach `energy_limit`. Absorbing boundaries (MUR, and especially PML_8) and structures with high-Q resonances both increase the number of time steps needed. Check the mesh cell count reported at the start of the run, and whether `refined_cellsize` is finer than actually needed everywhere, not just at the features that need it.

Can I reduce the stackup by removing the substrate for transmission-line models? Yes, if a ground plane between the signal and the substrate effectively shields the fields from reaching it — `SG13G2_nosub.xml` is provided for exactly this case, and lets you use PEC boundaries instead of needing absorbing ones for the substrate side.

Can I use a different data directory location? Yes — `sim_path` (built from `utilities.create_sim_path()`) defaults to `output/<model_basename>_data` next to your script, but you can point it anywhere by changing that code. If the model's own directory isn't writable, output falls back to a directory under your system's TEMP location automatically.

My imported GDSII looks strange. Check that you're reading the correct cell. The top-level cell is used by default; pass `cellname=` to `read_gds()` otherwise. Also check the correct `purpose` list, usually `[0]`. Holes and cutouts are handled automatically by the GDSII reader, see [Input files](#). If geometry still looks wrong, check for self-intersecting or degenerate polygons in the source GDSII itself.

My results show a lot of ripple. This is usually caused by one of two things. Either the simulation time is insufficient: `energy_limit` is not low enough, so the time-domain signal hasn't fully decayed before the FFT is taken. Or there are reflections from an under-sized simulation boundary: too little `margin/air_around`, or a boundary condition that doesn't match the physics, for example a PEC box where an absorbing boundary was needed.

Where do I change the view style in AppCSXCAD's 3D viewer? Use the viewer's own menus (Tools > Visibility, and the object tree) to select which geometry groups are shown — this is the same viewer used throughout this guide's screenshots.

To see a 2D view with no perspective distortion, use the 2D view mode, and set the axis accordingly. You can also move the displayed mesh plane position: it can be moved in z direction to top, so that the mesh lines are not hidden by metals.




Figure 10: AppCSXCAD view in 2D mode

I spotted a model mistake in the 3D preview — can I quit without simulating? Yes — close the AppCSXCAD window, or press Ctrl-C in the terminal, and fix your model script; nothing is simulated until you close the preview (with `preview_only = False`) or don't request a preview at all (`no_gui = True`).

15 Appendix

15.1 Keeping the local-copy example in sync

`more_examples/local_modules_copy/modules/` is a real (not symlinked) copy of `workflow/modules/`, kept in sync with `scripts/sync_local_modules_copy.py` — see `scripts/README.md`. Run it whenever you change `workflow/modules/` and also maintain a local checkout of that example.

15.2 List of examples

See [Examples](#) above for the current list, and the top-level [README.md](#) for longer per-example descriptions and result screenshots.