IHP SG13G2 · OPEN SOURCE RFIC EM WORKFLOW

Using openEMS with IHP SG13

User's Guide

Volker Mühlhaus
volker@muehlhaus.com

Document version 2026-10-31 (v3)

Contents

1 What's New 3

2 About this workflow 3
2.1 Method of Moments / FEM vs. FDTD 3

3 Getting the code: pip package vs. local copy 4
3.1 Recommended: install gds2openEMS as a Python module 4
3.2 Alternative, no longer recommended: local modules copy 5

4 Required software and Python modules 5

5 Model configuration syntax: settings{} dict vs. loose variables 5
5.1 Converting older models to the settings syntax 5

6 Quick tour 6
6.1 Minimum configuration 6
6.2 Input files 7
6.3 Port definitions 7
6.4 Meshing 9
6.5 Running the model: preview, mesh, simulate 10
6.6 Re-running a model: the hash-based skip 11
6.7 Output files and Touchstone conversion 12

7 The settings{} dictionary in detail 12
7.1 Required settings 13
7.2 Optional settings 14

8 Simulation volume and boundaries 14

9 Field dumps 16

10 Antenna simulation 16

11 Port parasitics and de-embedding 18

12 Advanced topics 18
12.1 Forcing the solver thread count 18
12.2 Running on the GPU 18
12.3 Resonance estimation: stop when the result has converged 19
12.4 Automatic meshing with easyMesh4openEMS 19
12.5 Overriding XML stackup `<Variable>` from the model code 20

13 Examples 20

14 FAQ 21

15 Appendix 22
15.1 Keeping the local-copy example in sync 22
15.2 List of examples 22

1 What's New

This chapter gives a brief overview of major changes since the previous edition (v2, March 2025) of this guide. For the complete, dated change log, see `CHANGES.md` in the repository.

- One major change is that documentation and examples now **install gds2openEMS as a module from PyPi** (pip install gds2openEMS), instead of the local copy of the modules folder used before. The old method with local copy still works, but the new method using Python module installation is more convenient.
- The second major change is that examples and documentation now use the **settings[] dictionary syntax for simulation model settings**, similar to the gds2palace workflow. The previous model syntax with loose variables is still supported, so you can choose whatever suits your needs. Long term, the settings[] dictionary syntax is preferred because it allows to switch between openEMS and Palace workflows more easily.
- A graphical **XML stackup editor** was added in the `stackup_editor` directory. When the gds2openEMS Python module is installed, you can call this from the commandline using `stackupEditor`, with an XML file as optional parameter. Because of this GUI-based stackup editor, the PySide6 module is now a required dependency, and installed automatically when you install the gds2openEMS module.
- **Derived layers and `<Variable>` overrides** in the XML stackup format, see Overriding XML stackup `<Variable>` from the model code and `XML_stackup_format.md`.
- The `postprocess_only` setting is no longer required. `runSimulation()` already skips the FDTD solve automatically via a content hash of the model. If nothing changed since the last run, it reuses the existing result instead of re-simulating.
- **`shapely` is now a hard dependency.** It is used for splitting keyhole and hole polygons before CSXCAD extrusion. Layouts with cutouts no longer need a `preprocess_gds` flag at all.
- Port de-embedding script `scripts/deembed_openEMS.py` is a "quick & dirty" solution to estimate and remove lumped port's parasitic inductance, see Port parasitics and de-embedding.
- `numThreads` setting to **force the solver thread count** to a fixed value, instead of relying on automatic detection.
- **Resonance estimation** (`settings['resonance_estimation']`): openEMS stops when the extrapolated S-parameters have converged, instead of at a fixed `energy_limit`, see Resonance estimation.

2 About this workflow

openEMS is an open-source EM simulation tool based on the FDTD (Finite Difference Time Domain) method. This workflow generates openEMS simulation models directly from **GDSII layout files** and an **XML technology stackup**, for RFIC structures on the IHP SG13G2 process — inductors, transmission lines, couplers, antennas and similar passive structures, producing S-parameters.

Two files must be provided by the user: the **layout in GDSII format**, and a **simulation model script** in Python. The model script references an XML stackup file and calls functions from the workflow code to build the openEMS/CSXCAD model, mesh it, run the FDTD solve, and post-process the result into Touchstone S-parameters. That workflow code can be the `gds2openEMS` package, or a local `modules/` copy — see the next chapter.

2.1 Method of Moments / FEM vs. FDTD

It's useful to understand how FDTD's simulation flow differs from a frequency-domain solver. Frequency-domain solvers include Method of Moments, as used in Momentum or Sonnet, and FEM, as used in the sibling `gds2palace_ihp_sg13g2` workflow:

- **MoM/FEM** is a frequency-domain solver: it solves once per frequency point and gets the **full S-matrix** for that one frequency. Covering a frequency range means repeating the solve at every frequency of interest.
- **FDTD** is a time-domain solver: exciting one port with a wideband Gaussian pulse and running the field propagation forward in time gives, after an FFT, **one column of the S-matrix across the entire frequency range** in a single run. To get the full S-matrix, you repeat the run with excitation at each other port in turn.

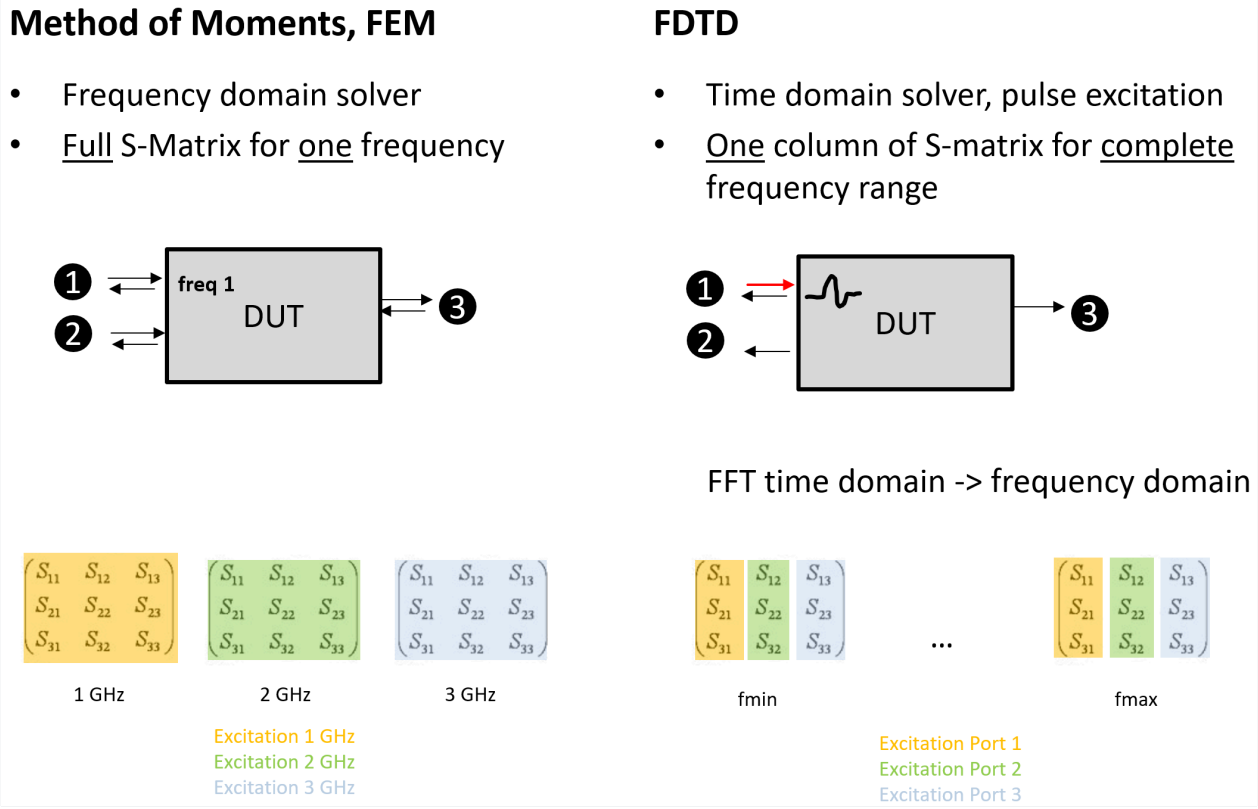


Figure 1: Comparison between MoM/FEM and FDTD simulation methods

The number of frequency points you request has essentially no effect on FDTD simulation time — the frequency-domain result is FFT post-processing of the time-domain data, so there's no speed penalty to specifying `numfreq = 401` instead of a coarser sweep.

3 Getting the code: pip package vs. local copy

There are two independent ways to get this workflow's code into your model script — pick one, they're interchangeable, and this choice is separate from the settings{} vs. loose-variable syntax choice below.

3.1 Recommended: install gds2openEMS as a Python module

```
python3 -m venv ~/venv/openems
source ~/venv/openems/bin/activate
pip install gds2openEMS
```

Your model script then does:

```
from gds2openEMS import *
```

and needs no local copy of the workflow code at all — just your GDSII layout, your XML stackup file, and the model script itself. Update with `pip install gds2openEMS --upgrade`. All examples under `workflow/` use this.

3.2 Alternative, no longer recommended: local modules copy

Clone/download this repository and copy the `workflow/modules/` folder next to your own model script, which then does:

```
import sys, os
sys.path.insert(0, os.path.abspath(os.path.join(os.path.dirname(__file__), 'modules')))
from modules import *
```

This is still fully supported — useful if you want to vendor a specific version, work offline, or edit the workflow code itself. `more_examples/local_modules_copy/` is a fully self-contained example written this way: its own `modules/` copy, GDSII layout and XML stackup, no PyPI package needed.

4 Required software and Python modules

This workflow is based on the Python bindings for openEMS; please refer to https://www.openems.de/ and https://docs.openems.de/python/install.html#python-linux-install for building/installing openEMS and CSXCAD themselves. Pre-built binaries are available for Windows, which is the easier path if you don't want to build from source; for Linux build issues, see the openEMS forum at https://github.com/thliebig/openEMS-Project/discussions.

In addition to openEMS/CSXCAD, these Python modules must be installed: `gdspy`, `shapely`, `numpy`, `matplotlib`. `pip install gds2openEMS` declares `gdspy`, `CSXCAD`, `openEMS`, `numpy` and `shapely` as its own dependencies, and pip will attempt to pull them in. Note that openEMS/CSXCAD's native bindings are normally built from source rather than installed as an ordinary wheel — see the openEMS install docs linked above. `matplotlib` is only needed by the example scripts' own plotting code, not by the installed package itself.

5 Model configuration syntax: settings{} dict vs. loose variables

Model scripts pass their configuration to `setupSimulation()`/`runSimulation()` two ways — pick one, the underlying functions accept both, and you can mix either syntax with either distribution method above:

- `settings = {}` **dictionary** is the recommended style. Every option is a key in one dict, passed as `setupSimulation(FDTD=FDTD, settings=settings)`. This is the style all `workflow/` examples use now, and it's also what `gds2palace_ihp_sg13g2` (the sibling AWS Palace/Elmer FEM workflow) uses. As a bonus, `max_cellsize` is computed for you from `fstop/unit/cells_per_wavelength` instead of you having to compute it yourself.
- **Loose top-level variables** (`fstart = 0e9`, `refined_cellsize = 1`, …) passed individually and positionally — the original style, still fully supported. `more_examples/local_modules_copy/run_line_viaport.py` is a current example of it.

The rest of this guide shows the `settings{}` style, since that's what all current `workflow/` examples use; anywhere it says `settings['xyz'] = ...`, the loose-variable style has an equivalent plain `xyz = ...` passed positionally instead.

5.1 Converting older models to the settings syntax

The tool `convert2settings` converts a model script from the loose-variable syntax to the `settings{}` syntax. It is installed together with gds2openEMS (version 0.6.5 or newer) and runs in a terminal:

1. Open a terminal (Windows: Command Prompt or PowerShell).

2. Activate the Python environment in which gds2openEMS is installed, the same one you use to run your models, e.g. `source ~/venv/openems/bin/activate` on Linux and macOS, or `C:\venv\openems\Scripts\activate` on Windows (use your own environment folder).
3. Go to the folder of the model and run:

```
convert2settings my_model.py
```

If the terminal reports that `convert2settings` is not found, the environment is not active, or its gds2openEMS is older than 0.6.5. `python -m gds2openEMS.convert2settings my_model.py` does the same and also works without the `Scripts` / `bin` folder in the search path.

To see the changes in a window first, with the original and the converted model side by side, add `--gui`:

```
convert2settings --gui my_model.py
```

It writes `my_model_settings.py` next to the original. The original is not changed, and an existing `my_model_settings.py` is not overwritten.

- Loose variables become `settings['...']` entries, and the `setupSimulation()` / `runSimulation()` calls get `settings=settings`.
- A model that imports a local `modules` folder is switched to the installed gds2openEMS package, and the `sys.path` line for that folder is removed. The converted model then uses the current workflow, not the old copy.
- The `max_cellsize` calculation is removed when it is the same as the one in `setupSimulation()`, which then does it.
- A variable is only moved into `settings[]` when its meaning is certain, because it is passed to a workflow function or used in one of the standard expressions such as `SetGaussExcite()`. File names and other variables stay as they are. A model that can't be converted safely is refused, with the reason.

Options:

Option	Effect
`--gui`	opens a window that shows the original and the converted model side by side, with the changed lines marked and each variable's reason, before anything is written
`--engine gpu`	adds `settings['engine'] = 'gpu'`, see Running on the GPU
`--dry-run`	shows the changes, writes nothing
`--json`	prints a detailed report in JSON

Several files or whole folders can be given at once (`convert2settings folder` converts all `*.py` files in it). Check the converted model in the AppCSXCAD preview before the first run: because it uses the current workflow, its results can differ slightly from those of an old `modules` copy.

6 Quick tour

This chapter gives a brief, practical walkthrough. It uses `workflow/run_inductor_diffport.py` (a 1-port octagon inductor) as the running example.

6.1 Minimum configuration

If you installed `gds2openEMS` via pip, the minimum configuration for a model is just three files in one directory: the **XML stackup**, the **GDSII layout**, and your **model script**. No `modules/` folder needed.

If you're using the local-copy method instead, add a fourth item: your own copy of the `modules/` folder, next to the script.

6.2 Input files

```
gds_filename = "L_2n0_simplified.gds"   # geometries
XML_filename = "SG13G2.xml"             # stackup

settings['preprocess_gds'] = False      # see note below
settings['merge_polygon_size'] = 1.0    # merge via arrays within this distance (microns), 0 to disable

materials_list, dielectrics_list, metals_list = stackup_reader.read_substrate(XML_filename)
layernumbers = metals_list.getlayernumbers()
layernumbers.extend(simulation_ports.portlayers)

allpolygons = gds_reader.read_gds(gds_filename, layernumbers, purposelist=[0],
                                  metals_list=metals_list,
                                  preprocess=settings['preprocess_gds'],
                                  merge_polygon_size=settings['merge_polygon_size'])
```

`preprocess_gds` is a legacy setting kept only for backward compatibility with older model scripts. The GDSII reader now always splits keyhole/hole polygons safely after flattening the layout, so this option is a no-op today. Setting it to `True` just prints a note that it's obsolete. New model scripts don't need to set it at all.

`merge_polygon_size` applies to layers declared `Type="via"` in the XML stackup. With a non-zero value, polygons on that layer are oversized by half that distance, overlapping ones are merged, then undersized back — merging a via array into one bounding-box shape, as long as the maximum via spacing doesn't exceed the given value.

Merging fills the gaps between the vias with via material, so the merged block conducts better than the real via array. The via conductivity in the SG13G2 stackup files is derived from the per-via resistance in the process specification, i.e. it is valid for the via cross section only. With `settings['fill_factor_correction'] = True`, `setupSimulation()` corrects this: each merged via polygon gets its conductivity multiplied by its fill factor, the original via area inside it divided by the merged area (rounded to 2 decimals). Polygons with different fill factors get separate materials, named like `TopVia2_x0.49`, which you can see in AppCSXCAD; single vias that were not merged keep fill factor 1 and the original material. The correction needs `merge_polygon_size > 0`, because the fill factor is calculated from the via polygons before merging: a GDSII file where via arrays were already merged into blocks (e.g. by `gds_prepare_for_EM`) gets no correction. PEC via layers are not scaled.

6.3 Port definitions

Ports are created from polygons on special GDSII layers (by convention, layer 201 and above), not by coding their position directly in Python:

```
simulation_ports = simulation_setup.all_simulation_ports()

# via port: from_layername / to_layername, direction is the vertical axis
simulation_ports.add_port(simulation_setup.simulation_port(
    portnumber=1,
    voltage=1,
    port_Z0=50,
    source_layernum=201,
    from_layername='Metal1',
    to_layername='TopMetal2',
    direction='z'))

# in-plane port: target_layername, direction is the propagation axis in-plane
simulation_ports.add_port(simulation_setup.simulation_port(
    portnumber=1,
    voltage=1,
    port_Z0=50,
    source_layernum=201,
    target_layername='TopMetal1',
    direction='x'))
```

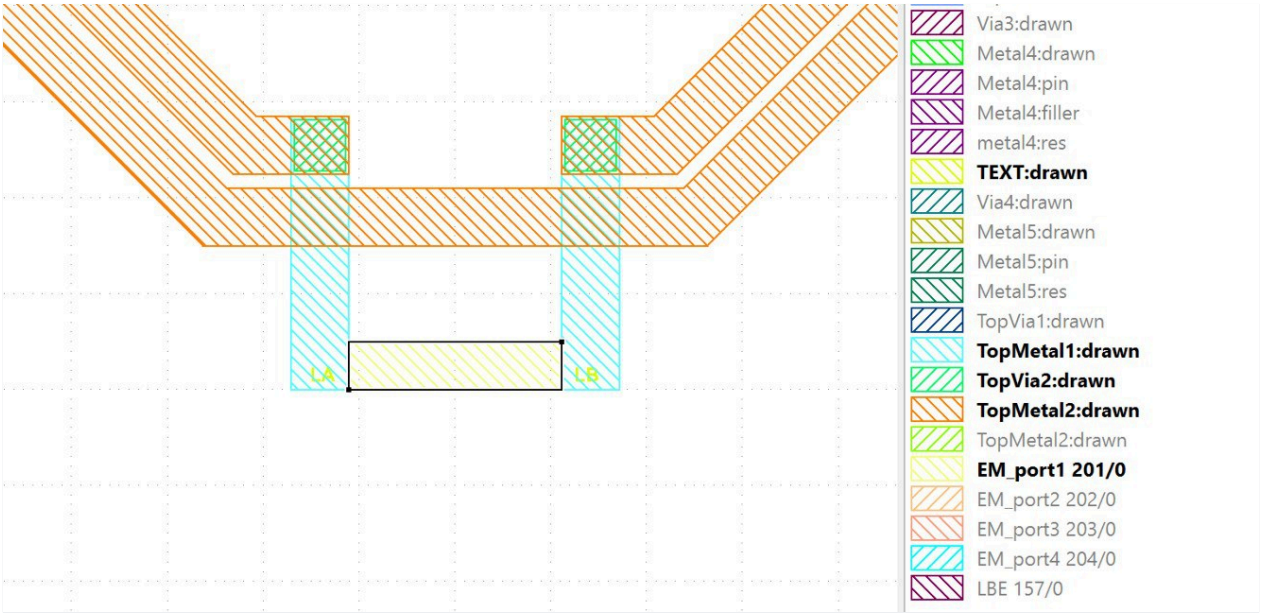


Figure 2: Port layout defined in GDSII file on layer 201

Each port must use its own GDSII source layer. `voltage=0` disables that port's own excitation without removing its definition.

In scripts that build their excitation list from `simulation_ports.all_active_excitations()` — like `run_generic_nport.py` — this also skips that port's FDTD run entirely, saving simulation time. Scripts with a hardcoded excitation list (most of the `workflow/` examples) must leave that port's number out of the list themselves; `voltage=0` alone does not skip anything there, since the script controls which ports run, independent of their `voltage` value.

A skipped port's S-parameters are simply not computed. If you then try to write a Touchstone file that references one of them, `write_snp()` raises a clear error naming the missing S-parameters, rather than producing an incomplete or corrupt file. See the FAQ.

The port can be mapped to a metal layer as shown here, with x or y orientation, but it can also be mapped to a via layer and used in z direction, if necessary. This gives maximum flexibility in creating EM ports.

The port direction also supports reverse polarity, by adding a minus sign to the direction. This can be required to place the reference terminal on the correct side, as shown in the example below. A via port with `'z'` runs from its lower to its upper metal, whatever the order of `from_layername` and `to_layername`; `'-z'` reverses it. Use the same orientation for all ports of a model, for example `'-z'` where the signal metal is below its reference metal.

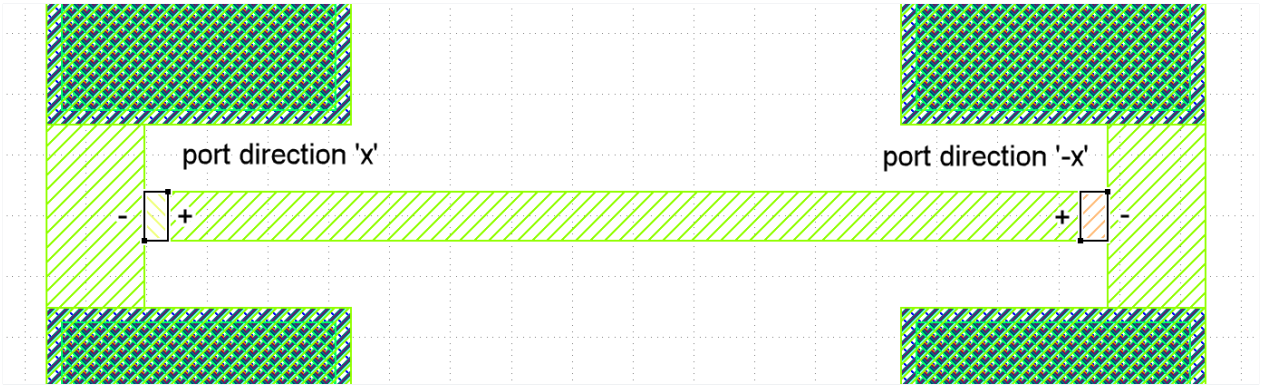


Figure 3: Example for polarity in port direction

Composite ports combine multiple EM ports into one logical port, for example a GSG pad pair. These are not built into `simulation_setup`'s port model directly.

The GSG port on each end of the line consists of 2 port definitions each, with one port in y direction and the other port in -y direction, to have the same effective polarity. And because they are effectively in parallel, we specify 2*Z0 as port impedance, to have Z0 in total.

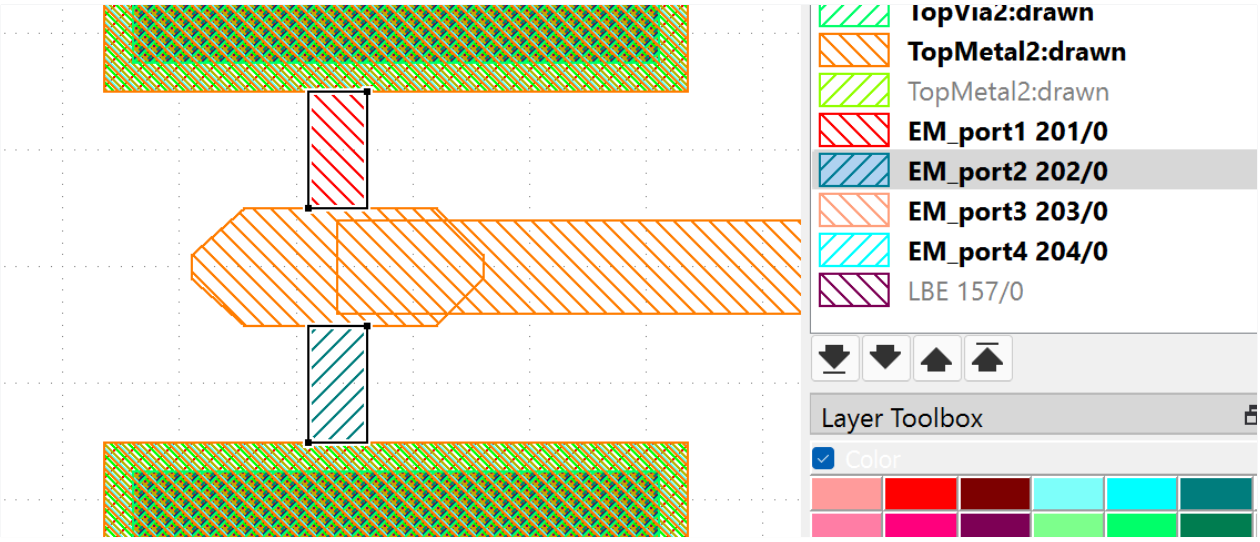


Figure 4: Port layout for GSG configuration, composite port from 2 port elements

`run_line_GSG_complex.py` shows the manual approach: simulate the underlying EM ports together, then combine their results in your own post-processing code.

Ports with several polygons. Several separate polygons on one port layer are the elements of one port, in parallel. Use this for a terminal that is spread over several strips, for example the source or drain fingers of a transistor. Each element is an openEMS lumped port with the resistance N·Z0 for N elements, so the elements in parallel terminate the port with Z0, and all elements are excited together. The port voltage is the mean of the element voltages, the port current the sum of the element currents; the S-parameters are calculated from these with the port's Z0, as for a single polygon.

Polygons on one port layer that touch or overlap are one element, with their common bounding box: two lumped ports on a shared edge would both act on that edge. setupSimulation() reports when it combines polygons.

Checked with `run_line_viaport.py`'s line, its ports drawn as one polygon or as two polygons with a 1 µm gap, on the same mesh: the S-parameters agree within |ΔS| 0.003 from 0 to 110 GHz, also with resonance estimation. Before gds2openEMS 0.6.7, every polygon became a port of its own with the same number: openEMS stopped with an error about a duplicate probe name, or, in older versions, ran with the port terminated by Z0/N.

6.4 Meshing

The default in all examples is **automatic meshing based on geometry**: mesh lines are placed to follow polygon edges and diagonal features. Lines that end up too close together are then merged or removed, since they would slow simulation without adding accuracy. `settings['refined_cellsize']` sets the target mesh resolution at conductor edges. `settings['cells_per_wavelength']` sets the coarse mesh size elsewhere in the model: the maximum cell size is the wavelength at `fstop` in the densest dielectric, divided by `cells_per_wavelength`.

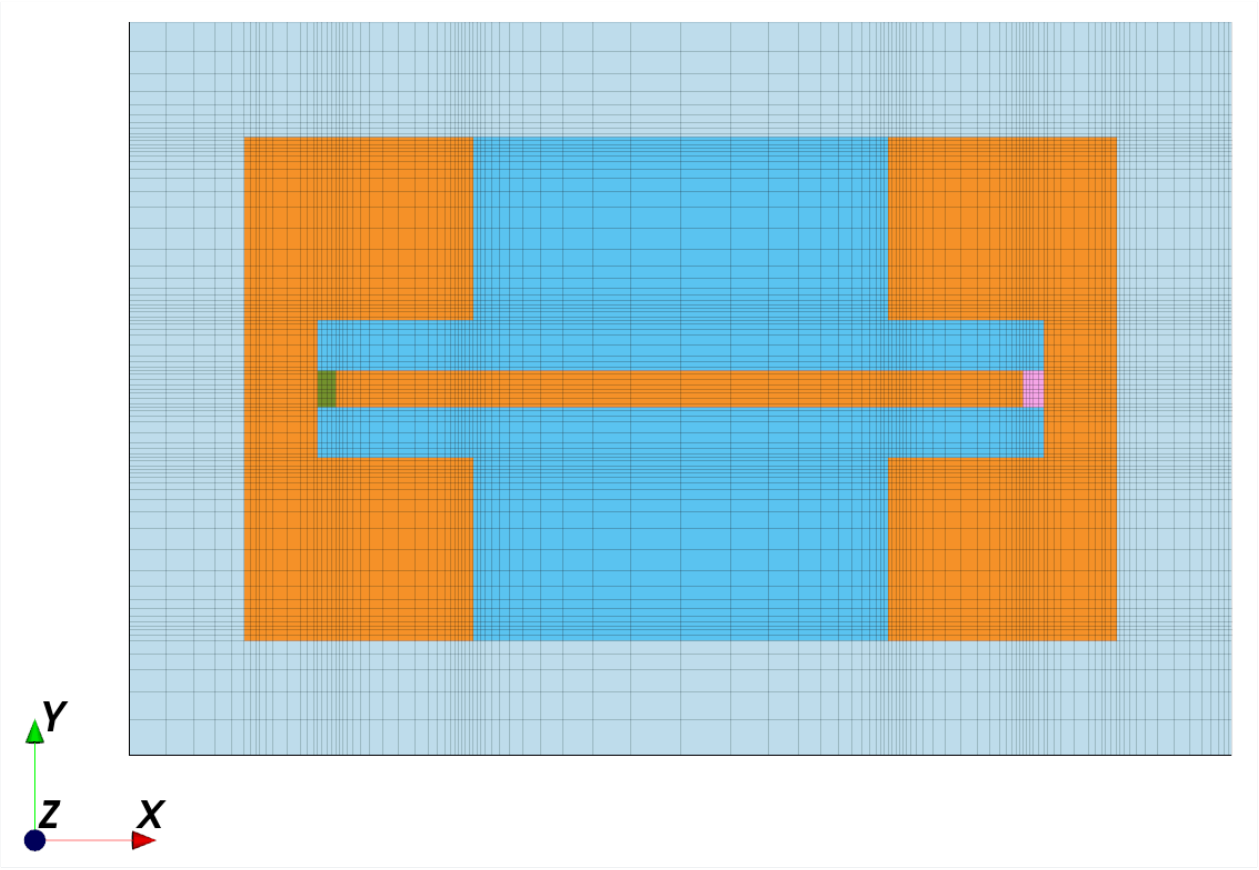


Figure 5: Automatic meshing

In z direction, mesh lines come from the stackup. Every dielectric boundary and the bottom and top of every metal and via layer used in the layout get a mesh line, and lines closer than 1 nm are merged into one. Metals are then subdivided:

- a metal up to 3× `refined_cellsize` thick is divided into cells of about `refined_cellsize`, split at any dielectric interface inside it (e.g. a passivation that cuts TopMetal2 halfway), so no subdivision line lands next to such an interface;
- a thicker metal gets one extra line at `refined_cellsize` from its bottom and top surface.

Dielectrics are graded between these lines, up to the maximum cell size. So the smallest z cell is set by the thinnest layer actually used, not by the meshing rules.

Before 28-Sep-2026, the z mesh placed metal subdivision lines without regard to interfaces inside the metal, which created very small cells in some stackups (0.1 µm in the L6n2 passivation-cut stackup at 1 µm `refined_cellsize`). To reproduce results from that version, set `settings['z_mesh_function'] = util_meshlines.create_z_mesh_legacy`.

The default meshing method described above is reliable and easy to work with, and is a good choice for most models. The optional example `more_examples/easyMesh/` demonstrates an alternative meshing engine, shown for information — in most cases the two approaches are equivalent, and neither is a general replacement for the other. See chapter Automatic meshing with easyMesh4openEMS.

6.5 Running the model: preview, mesh, simulate

```
FDTD = simulation_setup.setupSimulation(FDTD=FDTD, settings=settings)
sub1_data_path = simulation_setup.runSimulation(FDTD=FDTD, settings=settings)
```

With `settings['preview_only'] = True` and `settings['no_gui']` left at its default (`False`), running the script pops up **AppCSXCAD**, the 3D viewer, showing the meshed model before any simulation starts:

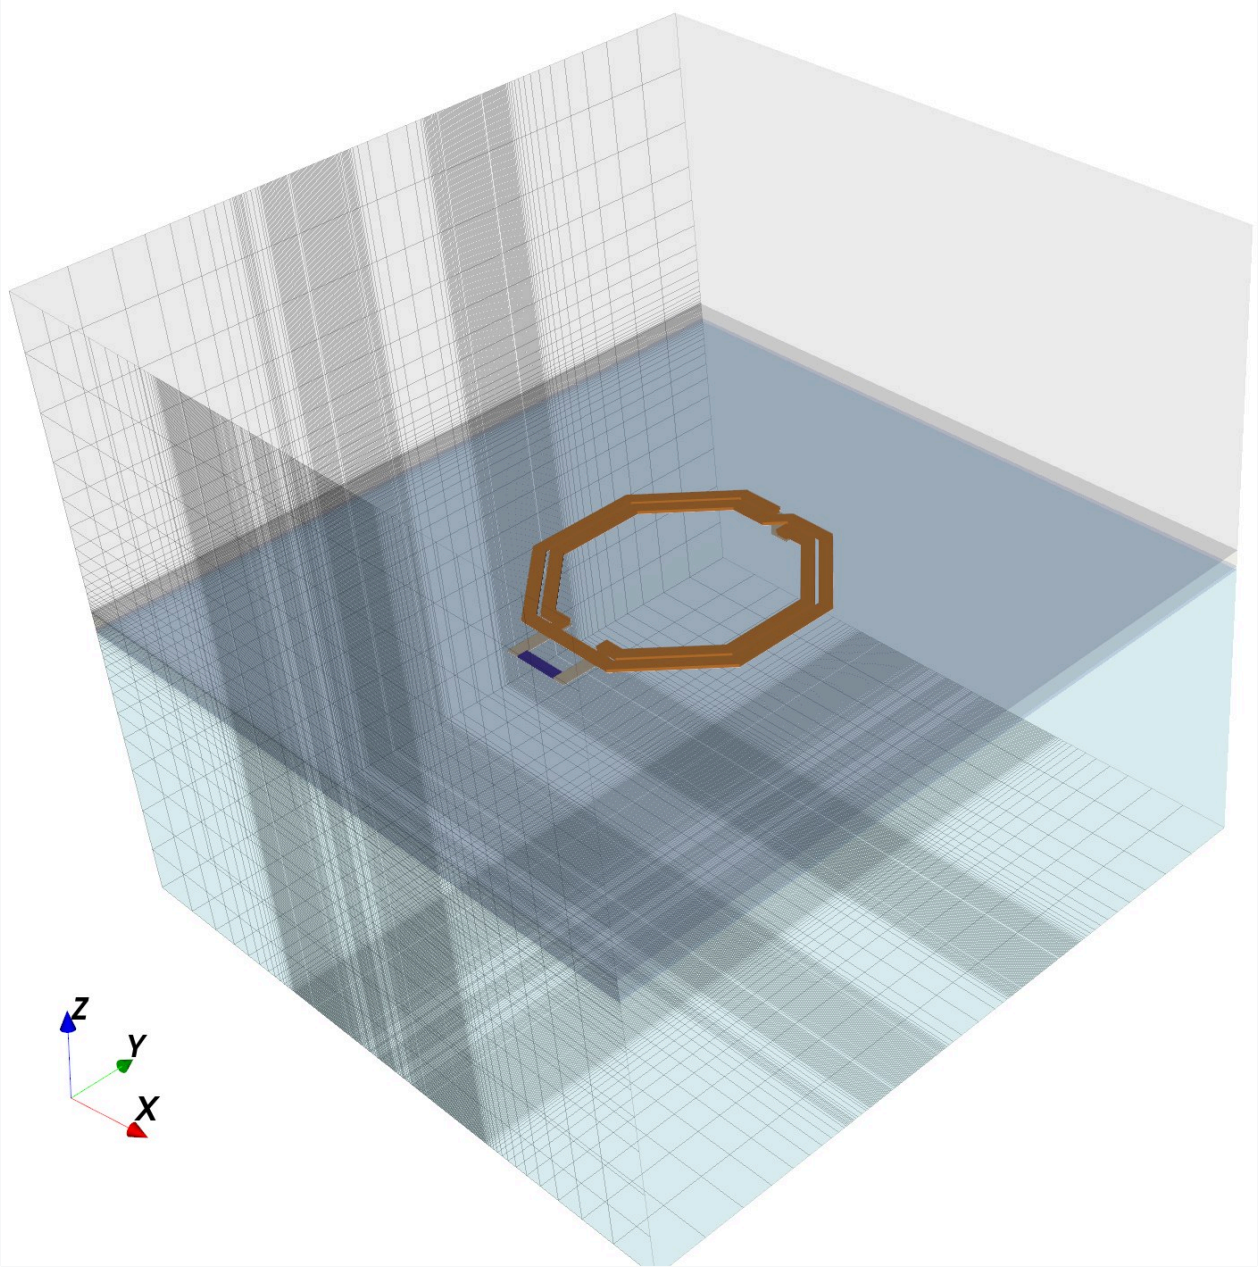


Figure 6: AppCSXCAD 3D mesh view of a differential-port octagon inductor

To close the preview, just close that window. You can also press Ctrl-C to abort. No other user action is required. Once you're happy with the model and mesh, set `settings['preview_only'] = False` and re-run. The FDTD solve then starts as soon as you close the preview window. If `settings['no_gui'] = True`, the solve starts immediately, since there is no preview window to close.

6.6 Re-running a model: the hash-based skip

`runSimulation()` computes a content hash covering both the generated CSX model file and the calling model script (`.py`) itself, up to and including the line that calls `runSimulation()`. Code further down in that script (post-processing, plotting, …) is not part of the hash. It compares that hash against the hash stored from the previous run in that model's output directory. If the hashes match, it skips the FDTD solve and prints:

```
Data for this model already exists, skipping simulation!
To force re-simulation, add parameter "force_simulation=True" to the runSimulation() call.
```

To force a re-solve even when the hashes match, pass `force_simulation=True` to `runSimulation()`.

The AppCSXCAD preview itself is **not** gated by this hash check. It is shown on every run, for the first port excitation, regardless of whether the model changed. Set `no_gui = True` to suppress it. This way you always get a chance to look at the current model before deciding whether to let it re-simulate. Only the FDTD solve is skipped when nothing changed. So you can freely re-run a script after only changing your own post-processing or plotting code, without waiting through a re-solve, while still seeing the preview every time.

6.7 Output files and Touchstone conversion

Output goes to `output/<model_basename>_data/sub-<n>/` for each port excitation `n`. This folder contains the CSX/mesh XML and the time-domain port voltage/current data. Field dumps are written there too, if you requested them.

There is no separate conversion tool for openEMS. Every example model script computes the S-parameters itself after the FDTD solve finishes, and calls `utilities.write_snp()` to write them straight to a Touchstone `.sNp` file:

```
snp_name = os.path.join(sim_path, model_basename + '.s' + str(num_ports) + 'p')
utilities.write_snp(s_params, f, snp_name, z0=simulation_ports.get_reference_impedance())
```

So the `.sNp` file already exists once your model script finishes running. Just run the script:

```
python your_model.py       # generates the model, runs the simulation, writes the .sNp file
```

7 The settings{} dictionary in detail

Settings keys are case-insensitive: `settings['numthreads']` works like `settings['numThreads']`, and `setupSimulation()` prints a note when it uses a key in a different case. If both spellings are set with different values, the spelling shown in the tables below wins, with a warning.

7.1 Required settings

Key	Meaning
`settings['unit']`	Unit of geometry values, typically `1e-6` (microns)
`settings['margin']`	Oversize of dielectrics from the GDSII bounding box, in the geometry unit
`settings['fstart']` / `settings['fstop']`	Start/stop frequency in Hz for the S-parameter output (FFT post-processing — see About this workflow
`settings['numfreq']`	Number of frequency points in the output sweep
`settings['refined_cellsize']`	Target mesh size at conductor edges
`settings['Boundaries']`	Required, no built-in default. List of 6 boundary conditions, one per side (`xmin, xmax, ymin, ymax, zmin, zmax`): `'PEC'` (lossless metal box), `'PMC'` (magnetic wall, useful for symmetry), `'MUR'` (simple absorbing), or `'PML_8'` (higher-quality absorbing, much slower simulation)
`settings['energy_limit']`	Residual energy (dB) at which the FDTD time-domain solve is considered converged
`settings['cells_per_wavelength']`	Coarse-mesh resolution away from refined edges, 10 or more (the examples use 10-20). The maximum cell size is calculated from it as wavelength at `fstop` / ($\sqrt{\varepsilon_{r,max}}$ · `cells_per_wavelength`). Alternatively, leave out `cells_per_wavelength` and set `settings['max_cellsize']` directly, in the geometry unit

7.2 Optional settings

Key	Default	Meaning
`settings['preview_only']`	`False`	Show the AppCSXCAD preview and stop, without simulating
`settings['no_gui']`	`False`	Never show AppCSXCAD; always proceed straight to simulation once the model changed (see the hash-based skip)
`settings['force_simulation']`	`False`	Re-simulate even if the model hash matches a previous run
`settings['max_cellsize']`	calculated	Maximum mesh cell size in the geometry unit. Only used when `fstop`, `unit` or `cells_per_wavelength` is missing; otherwise the value calculated from `cells_per_wavelength` is used
`settings['preprocess_gds']`	`False`	Legacy, now a no-op — see Input files
`settings['merge_polygon_size']`	`0`	Merge via-array polygons within this distance — see Input files
`settings['fill_factor_correction']`	`False`	Scale the conductivity of merged via arrays by their via fill factor — see Input files
`settings['air_around']`	`0`	Extra air spacing around the model in addition to `margin`, single value or a 6-element list `[xmin, xmax, ymin, ymax, zmin, zmax]`
`settings['numThreads']`	automatic	Force the openEMS solver thread count — see Forcing the solver thread count
`settings['engine']`	openEMS default	openEMS engine, e.g. `'gpu'` — see Running on the GPU
`settings['resonance_estimation']`	`False`	Stop openEMS when the extrapolated S-parameters have converged; `energy_limit` stays the upper limit — see Resonance estimation
`settings['easyMesh']`	`False`	Use the easyMesh4openEMS automatic mesh generator instead of the built-in one — see Automatic meshing with easyMesh4openEMS
`settings['field_dumps']`	none	A `simulation_setup.all_field_dumps()` object — see Field dumps

8 Simulation volume and boundaries

The simulation volume is the GDSII geometry bounding box, oversized by `margin` on all sides. If `air_around` is set, that adds further oversize. The volume is then meshed into cells according to `refined_cellsize` and `cells_per_wavelength`.

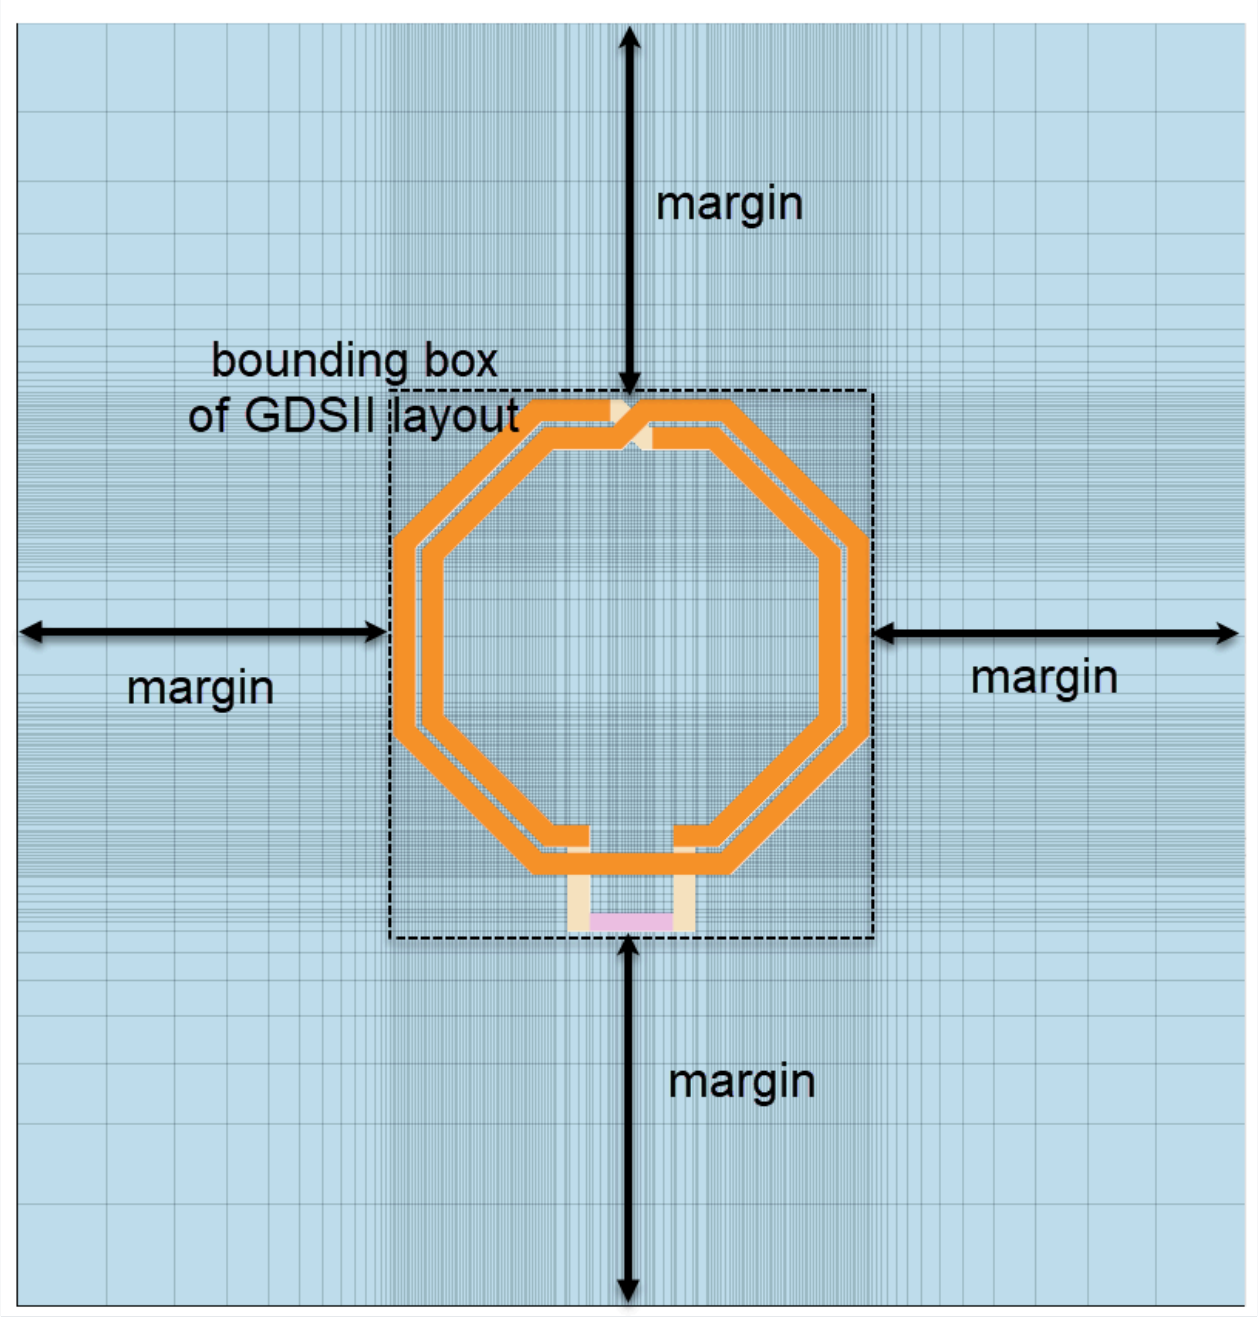


Figure 7: Margin defined in simulation settings

In the model code, boundaries are listed in this order: `xmin, xmax, ymin, ymax, zmin, zmax`. Each of the six outer boundaries is independently metal (`PEC`), magnetic-wall (`PMC`), or absorbing. Absorbing boundaries are `MUR`, or the more accurate but more expensive `PML_8`. See `settings['Boundaries']` above. A `PMC` boundary is useful to exploit a structure's symmetry and simulate only half the model. An absorbing boundary is needed wherever radiation should leave the simulation domain rather than reflect back in — for example all six sides for an antenna, see Antenna simulation.

Metal layers are often embedded inside a dielectric, so two materials would otherwise overlap. The workflow resolves this by construction: dielectric layers are the base fill, and metal/via geometry is meshed as its own conductor sheets or volumes on top, matching the Z-ranges defined for each layer in the XML stackup file.

9 Field dumps

For visualization in AppCSXCAD or ParaView, you can request field data to be written during simulation:

```
field_dumps = simulation_setup.all_field_dumps()
field_dumps.add_frequency_dump(name='Jf',
    file_type='vtk',
    dump_type='J',
    frequency=30e9,
    source_layernum=301,
    from_layername='TopMetal2',
    to_layername='TopMetal2',
    offset_top=0,
    offset_bottom=0)
field_dumps.add_time_dump(name='Et',
    file_type='vtk',
    dump_type='E',
    source_layernum=302,
    from_layername='Metal1',
    to_layername='TopMetal2',
    offset_top=10,
    offset_bottom=0)
settings['field_dumps'] = field_dumps
```

Each dump is positioned by its own GDSII bounding-box layer (like a port), with `dump_type` one of `'E'`, `'H'`, `'J'`, `'rotH'`, written as `'vtk'` or `'hdf5'`. That layer must be read from the GDSII file, like the port layers: add `layernumbers.extend(field_dumps.dumplayers)` before `read_gds()`. Otherwise the dump silently covers the bounding box of the entire layout, and `setupSimulation()` prints a warning. The dump box is the bounding rectangle of all polygons on that layer. Its edges also become mesh lines, like any other polygon edge. A frequency dump captures the steady-state field at one frequency; a time dump captures the field evolving during the FDTD solve. See `workflow/run_line_viaport_fielddump.py` for a complete example.

10 Antenna simulation

`workflow/run_dual_dipole.py` demonstrates far-field antenna pattern calculation (based on a 245 GHz dual-dipole design by IHP's Klaus Schmalz et al.). Antenna models differ from the passive-structure examples in three ways:

- All six outer boundaries are absorbing, either `MUR` or `PML_8`, since radiated power must leave the simulation domain rather than reflect back. Using an `'MUR'` boundary with a somewhat larger `air_around`/`margin` is usually faster than a tighter `'PML_8'` box for antenna problems. PML boundaries slow down openEMS a lot!
- In the simulation setup for this antenna, the optional parameter `air_around` adds additional air spacing all around the model, to ensure proper distance from the absorbing boundaries.
- An **NF2FF (near-field to far-field) box** is added around the radiator, which openEMS uses after the time-domain solve to compute the radiation pattern:

```
nf2ff_box = FDTD.CreateNF2FFBox(opt_resolution=[max_cellsize]*3, frequency=[ftarget])
...
nf2ff_res = nf2ff_box.CalcNF2FF(sub1_data_path, ftarget, theta, phi)
```

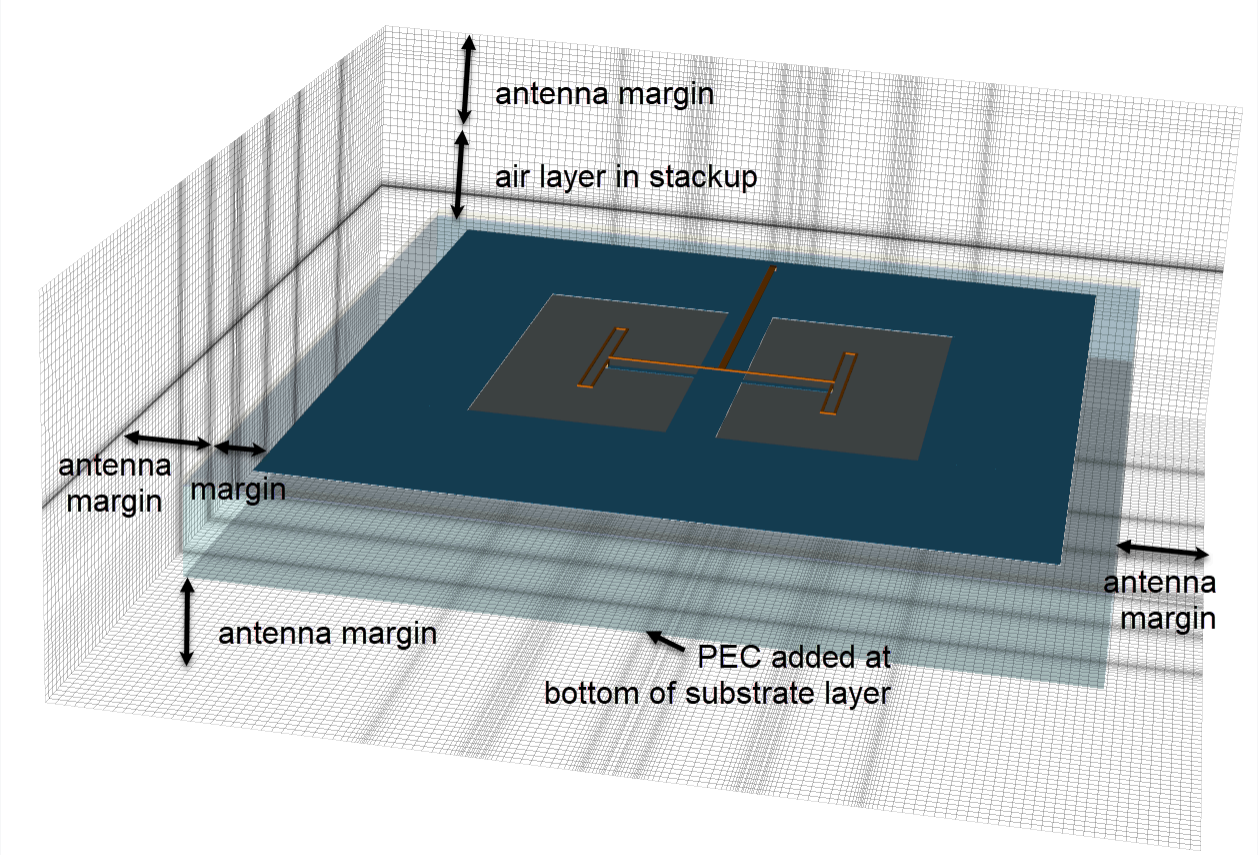


Figure 8: AppCSXCAD mesh view of the dual-dipole antenna model

From the NF2FF result and the port's accepted/incident power (`CSXport.P_acc/P_inc`), the model script computes total radiated power, peak directivity, radiation efficiency and mismatch loss, and plots directivity/gain vs. angle. See `run_dual_dipole.py`'s own code for the full calculation.

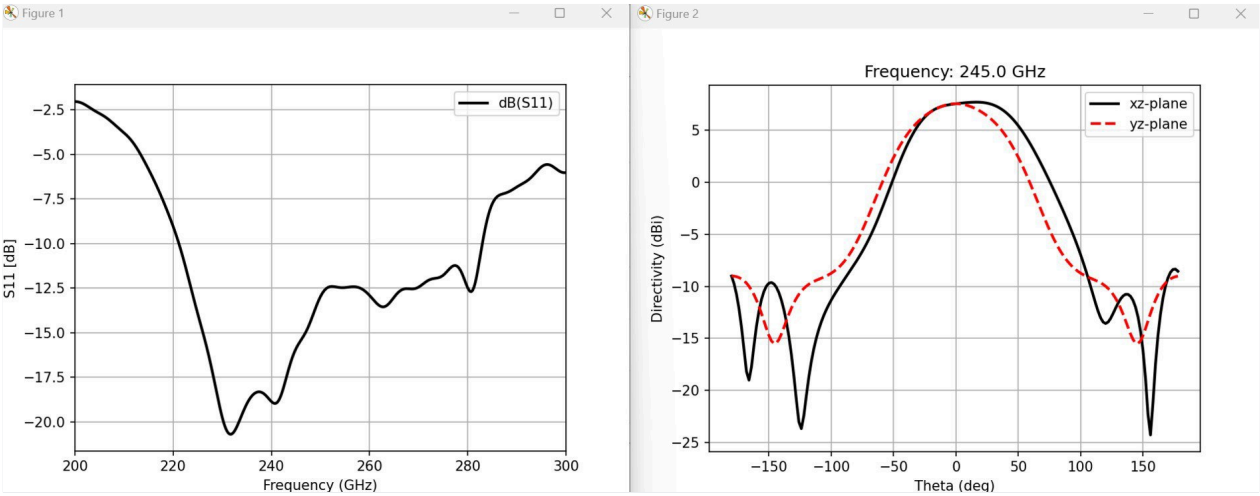


Figure 9: Return loss and directivity

This model needs `max_cellsize` for the NF2FF box resolution, and that value is needed *after* calling `setupSimulation()`. So it computes `wavelength_air/max_cellsize` manually, rather than relying on `setupSimulation()`'s internal settings-dict auto-computation. That computed value isn't written back into `settings` after the call.

The call to function `CreateNF2FFBox()` is plain openEMS syntax, so check out the documentation available there:

- https://docs.openems.de/python/openEMS/nf2ff.html

- https://github.com/thliebig/openEMS/blob/master/python/openEMS/nf2ff.py

Parameter `frequency` is optional, to store only the field data for that one frequency, which saves disk space. Without this optional parameter, nf2ff stores all time domain data, and patterns can be calculated in postprocessing for any frequency. The disadvantage is that data files will be much larger then.

For the evaluation of the NF2FF data files and creating the antenna pattern plot, we also use plain openEMS syntax. Have a look at example `run_dual_dipole.py` for more details.

11 Port parasitics and de-embedding

This workflow creates **lumped ports**, which introduce some physical length/area into the model and therefore some parasitic series inductance — port size that isn't small compared to the device under test will show up as extra inductance in the raw S-parameter result.

`scripts/deembed_openEMS.py` is a "quick & dirty" port de-embedding script: it reads the `port_information.json` geometry metadata that gds2openEMS writes automatically for each run, estimates the parasitic inductance built into each lumped port using a thin-sheet approximation, and removes it per port by cascading a negative-inductance section:

```
python3 deembed_openEMS.py inputfile.s2p
```

Output is written to a new file with suffix `_deembedded`. This works for any port count, but the method is approximate and **not applicable to composite port configurations** like `run_line_GSG_complex.py`'s GSG ports — treat it as a useful estimate, not a precise correction. For a port with several polygons, `port_information.json` has one entry per element, and the script takes the element inductances in parallel.

12 Advanced topics

12.1 Forcing the solver thread count

By default, openEMS auto-detects the "best" thread count, which can end up low if the machine is busy with other work. `settings['numThreads']` overrides that:

```
settings['numThreads'] = 8   # 0 (or omitted) = automatic detection
```

See `more_examples/numThreads/` for complete examples in both the `settings{}` and loose-variable styles.

12.2 Running on the GPU

openEMS has a GPU engine: Metal on macOS, HIP on AMD (ROCm) and NVIDIA (CUDA) GPUs. `settings['engine']` selects it:

```
settings['engine'] = 'gpu'   # omitted: openEMS default (multithreaded CPU engine)
```

The value is passed to openEMS as the `--engine` option, so other engines (e.g. `'multithreaded'`) work too.

Getting openEMS with the GPU engine: the GPU engine is newer than the latest openEMS release (v0.37.0), and the release packages, including the Windows package, don't have it yet. Until the next release, build openEMS from its current source code, openEMS-Project, as described in the openEMS installation guide, with the Python bindings. On macOS, the Metal backend is built by default. The HIP backend for AMD and NVIDIA GPUs needs ROCm, or CUDA with the HIP headers, on the build machine; see the "Build" section of the openEMS change log. In a build with the GPU engine, `openEMS --help` lists `gpu` for `--engine`.

- openEMS without the GPU engine (v0.37.0 and older) ignores the setting without an error and runs on the CPU. gds2openEMS prints a warning when it finds such an openEMS. The openEMS output shows which engine runs: `Create FDTD engine (GPU, backend: ...)` for the GPU engine.
- openEMS with the GPU engine stops with an error if it finds no usable GPU device, and gds2openEMS stops too. It doesn't fall back to the CPU.

Changing the engine is a change to the model script, so the simulation runs again instead of reusing existing results.

With PML boundaries the GPU engine gains most, because the CPU engine is slow for PML. On an Apple M4 (Metal), `run_line_full2port.py` (PML) ran 3.5 times faster than on the CPU, `run_line_GSG_complex.py` (PEC and MUR boundaries) 1.5 times. The S-parameters differ only slightly, mostly because the run ends at a slightly different timestep. If a part of the model has no GPU implementation, openEMS prints a warning listing it; it then runs on the CPU, and the speed gain is mostly lost.

12.3 Resonance estimation: stop when the result has converged

With `energy_limit`, openEMS stops when the energy in the model has decayed by that amount. The right value depends on the model. At −40 dB, a slow decay can be cut off, for example the L/R tail of an inductor or a time constant of the substrate. The S-parameters then get a truncation error, largest at low frequency: for an inductor, the low-frequency series resistance can be 25 % too low. At −60 dB, many models run longer than they need.

`settings['resonance_estimation'] = True` checks the result instead of the energy. While openEMS runs, the port voltages and currents are extrapolated beyond the current time, similar to the "AR filter" in CST or "resonance estimation" in Empire XPU, and openEMS is stopped as soon as the extrapolated S-parameters no longer change. The S-parameters are then calculated from the extrapolated signals.

```
settings['energy_limit'] = -60            # upper limit: openEMS stops here at the latest
settings['resonance_estimation'] = True   # stop earlier, once the extrapolated S-parameters have converged
```

In tests with six models (three inductors, a MIM capacitor, two PA core layouts up to 350 GHz), the runs stopped after 0.45–0.85× the time of a −60 dB run. Against a −90 dB reference run, the S-parameters differed by at most $9{\cdot}10^{-5}$ to $1.5{\cdot}10^{-3}$ (max $|\Delta S|$), depending on the model. For the inductors and the MIM capacitor, this was as accurate as or more accurate than a plain −60 dB run. The example `more_examples/resonance_estimation` shows the problem and the results step by step.

- `energy_limit` stays the upper limit. If the extrapolation has not converged by then, openEMS stops there as usual, and the port signals are still extended afterwards if the result passes a self-check.
- When resonance estimation stops openEMS, openEMS reports "Max. number of timesteps was reached before the end-criteria … was reached". That refers to this stop; no timestep limit was hit. The log file (below) states the reason.
- Each port excitation writes `resonance_estimation.txt` into its `sub-N` folder: why openEMS stopped, which fit was used, the slowest time constant found, and whether the extended signals are used. The extended signals are in `sub-N/resonance_estimation/`; the original openEMS data is not changed. `utilities.calculate_Sij()`, and the Y and Z functions based on it, use the extended signals automatically.
- With field dumps or nf2ff boxes in the model, openEMS is not stopped early, because these collect their data during the run and would be cut off. The port signals are still extended after the run.
- Changing the setting changes the model hash, so the model is simulated again (see the hash-based skip).

12.4 Automatic meshing with easyMesh4openEMS

`more_examples/easyMesh/openEMS_generic_nport_MA.py` demonstrates easyMesh4openEMS, an alternative automatic mesh-line placement engine, enabled via:

```
settings['easyMesh'] = True
```

`easyMesh4openEMS` is a separate package (`pip install easyMesh4openEMS`) not installed by default alongside `gds2openEMS`. If it isn't installed, `settings['easyMesh'] = True` doesn't crash the run — `setupSimulation()` catches the missing import, prints a note that it's being ignored until you install the package, and falls back to the default geometry-based automatic meshing.

12.5 Overriding XML stackup `<Variable>` from the model code

In the early days of this workflow, any change to the technology stackup like thickness of bulk silicon or thickness of air layer above required a separate, modified XML file.

To simplify this, an extension to the stackup file format was introduced with variable in the XML stackup file. You can now override these variables from your model code, to control values like thickness without using another XML file.

`more_examples/parameterized_XML_stackup/` shows how to override stackup variables `total_thickness`, `air_thickness` from the model script instead of editing the XML file directly. This is done in the `read_substrate()` code line:

```
materials_list, dielectrics_list, metals_list = stackup_reader.read_substrate(
    XML_filename, variable_overrides={'total_thickness': 100, 'air_thickness': 100})
```

Two versions of the same model are provided — `run_rfcmim_2port_parameters.py` (loose variables) and `run_rfcmim_2port_parameters_settings.py` (`settings{}` dict) — producing an identical model, to show both syntax options side by side. See the XML stackup's <Variables> block and `"="`-prefixed expressions in `XML_stackup_format.md` for the full format.

13 Examples

All examples read GDSII + XML stackup and write a Touchstone S-parameter file (except where noted). For a longer description of each, see the top-level README.md.

`workflow/` (primary examples, `settings{}` + `gds2openEMS`):

- `run_line_viaport.py` — thru line with via ports, single-side excitation with symmetry assumed for the reverse path
- `run_line_full2port.py` — thru line, full 2-port excitation
- `run_line_GSG_complex.py` — thru line with composite GSG ports
- `run_line_noGDSII.py` — geometry built entirely from Python code, no GDSII input
- `run_line_viaport_fielddump.py` — via-port line with field dumps (see Field dumps)
- `run_inductor_diffport.py` — 1-port octagon inductor (differential mode)
- `run_inductor_2port.py` — 2-port version of the same inductor, via ports to a common ground
- `run_generic_nport.py` — generic n-port model, excites every defined port in turn for a full S-matrix
- `run_rfcmim_2port_full.py` — MIM capacitor with via-array merging
- `run_dual_dipole.py` — antenna example, see Antenna simulation

`more_examples/`:

- `local_modules_copy/run_line_viaport.py` — the local-copy + loose-variable distribution/syntax combination, self-contained (own `modules/`, GDS, XML)
- `parameterized_XML_stackup/` — XML <Variable> overrides, both syntax styles
- `numThreads/` — forcing solver thread count, both syntax styles
- `easyMesh/openEMS_generic_nport_MA.py` — easyMesh4openEMS automatic meshing
- `resistors_sg13g2/run_resistors_Rsil.py` — IHP resistor derived-layer support
- `combine_layout_sources/run_combine_example.py` — combining GDSII-sourced and manually-coded geometry (rectangles/polygons added directly in Python) in one model

14 FAQ

How do I use the example models? Copy the example script, its GDSII file and XML stackup into a new directory, or into your own project directory. If you're using `gds2openEMS` via pip, no `modules/` copy is needed there. Then modify the configuration to match your own layout and requirements.

How do I create my own model? Start from the example that's closest to your use case, rather than starting from scratch. Match the port type, via or in-plane, and the port count. Then adapt the GDSII/XML filenames, port definitions and `settings{}` values.

Does `voltage=0` actually skip that port's simulation? It depends on how the script builds its excitation list. `run_generic_nport.py` and a few other examples call `simulation_ports.all_active_excitations()` to get the list of ports to excite, which already filters out `voltage=0` ports — for those scripts, a `voltage=0` port's FDTD run is skipped entirely. Most `workflow/` examples instead hardcode which port numbers to excite (e.g. `[1]` or `[[1],[2]]`); there, `voltage=0` by itself changes nothing — you must also remove that port number from the hardcoded list yourself. Either way, a port that was never excited has no S-parameters computed for it; trying to write a Touchstone file that references one raises a clear error naming the missing S-parameter(s), rather than producing a corrupt file.

What is a good mesh size? There's no single answer — it depends on the smallest feature (gap or line width) you need to resolve accurately, and on how much time you can afford. As a starting point, a `refined_cellsize` around 1/5 to 1/10 of your smallest critical dimension is reasonable; check convergence by comparing results at two different `refined_cellsize` values before trusting the finer one.

"Warning: Unused primitive (type: LinPoly) detected in property ...": These messages pop up when openEMS solver detects a layout with multiple touching metals on the same layer and same priority. It is purely cosmetic in this workflow and can be safely ignored.

Why is my simulation slow? FDTD simulation time scales with mesh cell count and the number of time steps needed to reach `energy_limit`. Absorbing boundaries (`MUR`, and especially `PML_8`) and structures with high-Q resonances both increase the number of time steps needed. Check the mesh cell count reported at the start of the run, and whether `refined_cellsize` is finer than actually needed everywhere, not just at the features that need it. Resonance estimation stops the run once the result has converged, instead of waiting for `energy_limit`.

Can I reduce the stackup by removing the substrate for transmission-line models? Yes, if a ground plane between the signal and the substrate effectively shields the fields from reaching it — `SG13G2_nosub.xml` is provided for exactly this case, and lets you use `PEC` boundaries instead of needing absorbing ones for the substrate side.

Can I use a different data directory location? Yes — `sim_path` (built from `utilities.create_sim_path()`) defaults to `output/<model_basename>_data` next to your script, but you can point it anywhere by changing that code. If the model's own directory isn't writable, output falls back to a directory under your system's TEMP location automatically.

My imported GDSII looks strange. Check that you're reading the correct cell. The top-level cell is used by default; pass `cellname=` to `read_gds()` otherwise. Also check the correct `purposelist`, usually `[0]`. Holes and cutouts are handled automatically by the GDSII reader, see Input files. If geometry still looks wrong, check for self-intersecting or degenerate polygons in the source GDSII itself.

My results show a lot of ripple. This is usually caused by one of two things. Either the simulation time is insufficient: `energy_limit` is not low enough, so the time-domain signal hasn't fully decayed before the FFT is taken (lower `energy_limit`, or use resonance estimation). Or there are reflections from an under-sized simulation boundary: too little `margin`/`air_around`, or a boundary condition that doesn't match the physics, for example a `PEC` box where an absorbing boundary was needed.

Where do I change the view style in AppCSXCAD's 3D viewer? Use the viewer's own menus (Tools > Visibility, and the object tree) to select which geometry groups are shown — this is the same viewer used throughout this guide's screenshots.

To see a 2D view with no perspective distortion, use the 2D view mode, and set the axis accordingly. You can also move the displayed mesh plane position: it can be moved in z direction to top, so that the mesh lines are not hidden by metals.

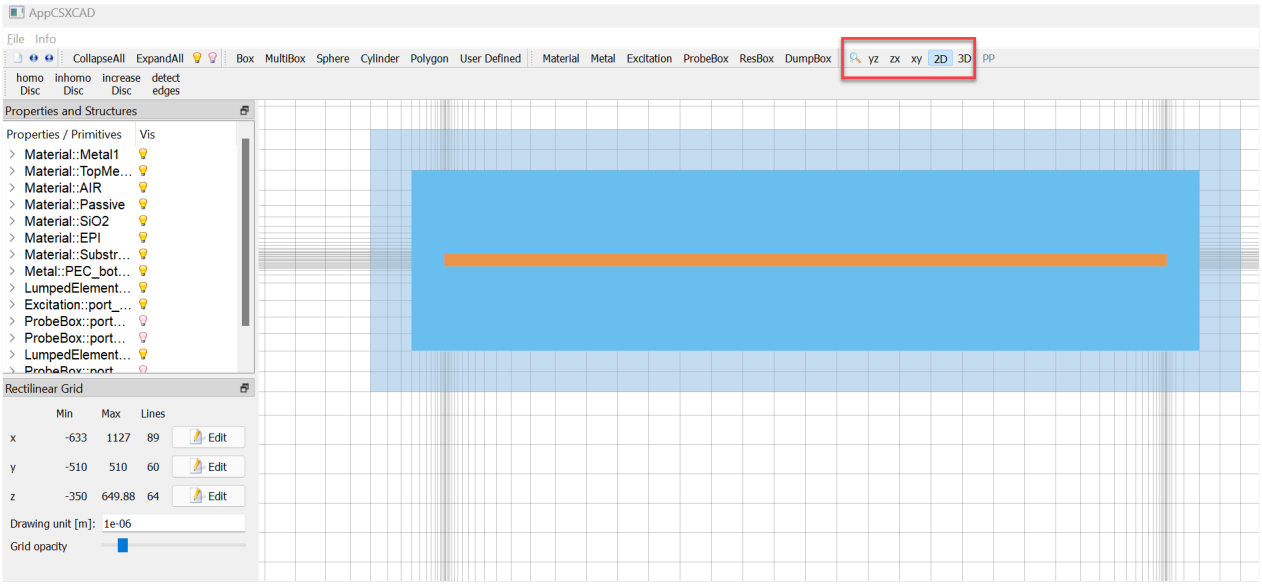


Figure 10: AppCSXCAD view in 2D mode

I spotted a model mistake in the 3D preview — can I quit without simulating? Yes — close the AppCSXCAD window, or press Ctrl-C in the terminal, and fix your model script; nothing is simulated until you close the preview (with `preview_only = False`) or don't request a preview at all (`no_gui = True`).

15 Appendix

15.1 Keeping the local-copy example in sync

`more_examples/local_modules_copy/modules/` is a real (not symlinked) copy of `workflow/modules/`, kept in sync with `scripts/sync_local_modules_copy.py` — see `scripts/README.md`. Run it whenever you change `workflow/modules/` and also maintain a local checkout of that example.

15.2 List of examples

See Examples above for the current list, and the top-level README.md for longer per-example descriptions and result screenshots.